Identifying Exoplanets with Deep Learning: A CNN and RNN Classifier for Kepler DR25 and Candidate Vetting

Bibin Thomas,[1,*] Vittal Bhat M,[2] Salman Arafath Mohammed,[3] Abdul Wase Mohammed,[4] Adis Abebaw Dessalegn,[5] and Mohit Mittal[6]

[1]*Centre for Research Impact & Outcome, Chitkara University Institute of Engineering and Technology*
[2]*Dept of ECE, P. A. College of Engineering*
[3]*Electrical Engineering Department, Computer Engineering Section, College of Engineering, King Khalid University, Abha, KSA*
[4]*Department of Electrical Engineering, College of Engineering, King Khalid University, Abha, KSA*
[5]*Department of Electrical and Computer Engineering, Faculty of Technology, Debre Markos University, P. O. Box 269, Debre Markos, Ethiopia*
[6]*Department of Data Science, Galgotias College of Engineering and Technology, Greater Noida, India*

ABSTRACT

The rapid expansion of exoplanet survey missions such as Kepler, TESS, and the upcoming PLATO mission has generated massive light-curve datasets that challenge traditional vetting pipelines. We introduce a hybrid deep-learning framework that integrates convolutional networks, bidirectional LSTMs, and an attention mechanism to identify planetary transit signals with improved accuracy and interpretability. Trained on Kepler DR25 data, the model achieves F1 = 0.910 ± 0.008 (AUC–ROC = 0.984 ± 0.004), significantly outperforming CNN-only baselines. Attention-based visualizations highlight ingress and egress phases consistent with astrophysical expectations. Applied to 1,360 DR25 candidate dispositions, our pipeline identified 190 high-confidence signals (P(model) > 0.70) that passed initial validation. Following comprehensive false positive probability (FPP) analysis and contamination tests, 13 candidates achieved statistical validation, of which three were fully confirmed as robust exoplanets: KOI-901.01 (warm mini-Neptune), KOI-1066.01 (hot Jupiter), and KOI-212.01 (warm Neptune). With an inference time of 80 ms per candidate, this framework enables scalable real-time triage for TESS and future PLATO operations, while the validated candidates provide a prioritized list for radial velocity, timing, and atmospheric follow-up. This interpretable, efficient pipeline addresses the bottleneck in exoplanet confirmation, advancing scalable discovery in the era of large-scale surveys.

1. INTRODUCTION

Over the preceding twenty years, there have been a lot of new discoveries of exoplanets made possible by large-scale photometric surveys like Kepler and TESS [7]. ESA's PLATO mission is set to launch in 2027 (European Space Agency 2025). These missions make light curves for hundreds of thousands of stars, which gives us a lot of data that can't be fully checked by individual or through conventional pipelines. TESS is currently on its second extended mission, which started in September 2022 and will last until 2027. This is to cover as much of the sky as possible and go back to fields it has already seen. TESS has found 7,655 candidate exoplanets as of July 2025, with 638 confirmed (NASA Exoplanet Archive Team 2025). This shows that we need scalable, automated ways to classify signals so we can tell which ones are real planets.

Machine learning methods are starting to be very important in this problem. A landmark study by (Shallue & Vanderburg 2018) demonstrated that convolutional neural networks (CNNs) could learn to distinguish real transits from false positives in Kepler data, significantly reducing the human workload in candidate vetting. Subsequent work has extended this idea to other archi-

---

Email: Bibinthomas951@gmail.com

* Scientist

[7] https://exoplanetarchive.ipac.caltech.edu/

tectures, including recurrent neural networks (RNNs), attention-based models, and hybrid frameworks. However, a key limitation of CNN-only methods is that they primarily capture local features in phase-folded light curves, while temporal dependencies and sequence-level context are often lost. This is especially critical for detecting weak or long-period planets, where subtle patterns across multiple transits determine the credibility of a candidate.

In this study, we hypothesize that explicit temporal modeling using bidirectional LSTMs (BiLSTMs), combined with an attention mechanism, can improve classification performance on ambiguous or low-SNR transit signals. The BiLSTM component enables the model to capture long-range sequential information across the light curve, while the attention layer provides interpretability by highlighting which parts of the sequence influence the classification decision. Together, these additions extend the (Shallue & Vanderburg 2018) paradigm, which we treat as our primary baseline.

We present a CNN–BiLSTM–Attention framework trained on Kepler DR25 light curves and applied to the DR25 catalog. Our model achieves improved performance over CNN-only baselines (F1 = 0.910 vs 0.879), and attention maps show that the network focuses on transit ingress/egress regions, providing interpretable insights into the classification process. When applied to 701 Kepler DR25 candidates, the model identified 136 high-confidence signals (P(model) > 0.70), from which 3 were validated as robust exoplanets following astrophysical vetting, with 13 additional candidates requiring follow-up observations.

Beyond Kepler, our framework has direct implications for ongoing and upcoming surveys. The methods developed here are immediately applicable to the TESS extended mission (through 2025), where systematic noise and limited baseline make candidate vetting challenging, and will be highly relevant for PLATO 's high-precision, long-baseline light curves. Furthermore, the prioritized candidates identified by our pipeline are well suited for atmospheric follow-up with the James Webb Space Telescope (JWST), which is already transforming exoplanet characterization.

In summary, this work contributes:

- A hybrid deep-learning framework that integrates CNN, BiLSTM, and attention for exoplanet detection.

- A demonstration of improved classification performance and interpretability compared to CNN-only models.

- Application of the model to Kepler DR25 candidates, yielding 3 validated exoplanets and 13 promising targets.

- A discussion of implications for current (TESS) and future (PLATO, JWST) missions.

This paper is structured as follows: Section 2 reviews related work; Section 3 describes our methodology; Section 4 outlines dataset construction; Sections 6–8 present training, experiments, and results; Sections 10–12 analyze validated candidates; and Section 15 concludes with future directions.

2. LITERATURE REVIEW

2.1. *Traditional Vetting Approaches*

Early exoplanet transit searches relied on classical algorithms such as the Box Least Squares (BLS) method[8] (Kovács et al. 2002) to detect periodic flux dips, followed by manual vetting by experts. With the large Kepler data releases, semi-automated tools like the Kepler Robovetter[9] (Coughlin et al. 2016) were introduced to filter astrophysical false positives such as eclipsing binaries, instrumental noise, and centroid shifts. While effective, these approaches struggled with scalability and were sensitive to noise artifacts, motivating the adoption of machine learning.

2.2. *CNN-Based Approaches*

A breakthrough came with (Shallue & Vanderburg 2018), who demonstrated that convolutional neural networks could directly classify Kepler light curves into planets and false positives. Their AstroNet architecture achieved state-of-the-art precision and enabled the discovery of new exoplanets from archival Kepler data. Subsequent work extended CNN-based approaches to TESS (Pearson et al. 2018), showing that automated deep learning classifiers could operate effectively even with shorter baselines and higher noise. CNNs remain attractive due to their ability to extract local transit-like features automatically, but their receptive field is limited to localized patterns rather than global temporal context.

---

[8] Box Least Squares (BLS) is a classical period-finding algorithm that searches for periodic, box-shaped dips in stellar light curves characteristic of planetary transits.

[9] The Kepler Robovetter is an automated classification system that applies multiple diagnostic tests to distinguish genuine planetary transits from astrophysical false positives in the Kepler candidate catalog.

2.3. Beyond CNNs: RNNs, Transformers, and Attention

To capture temporal dependencies, several works introduced recurrent neural networks (RNNs) alongside CNNs. For example, Pearson et al. (2018) combined CNN feature extraction with LSTMs to model transit sequences across multiple epochs. More recently, attention mechanisms and Transformer-based architectures have been applied to astronomical time series. (Yu et al. 2019) showed that self-attention can improve sensitivity to long-period and low-SNR signals, while introducing exoplanet-focused Transformers capable of handling irregular sampling and noisy baselines. These advances reflect a broader trend in time-series ML: combining convolutional, recurrent, and attention mechanisms to balance local feature extraction, sequential context, and interpretability. Despite these developments, challenges remain. RNN-based models capture long-range dependencies but often lack transparency, while Transformer models are data-hungry and computationally intensive. Importantly, very few studies have systematically explored hybrid models that integrate CNNs, BiLSTMs, and attention mechanisms specifically for exoplanet detection at scale. In the broader astronomical context, recent reviews have surveyed the application of machine-learning techniques to time-series data, framing exoplanet transit classification alongside other variable-star and transient science challenges (Baron 2019).

2.4. Gap and Our Contribution

This gap motivates our approach. CNNs alone cannot fully capture sequential context; RNNs address this but struggle with interpretability; attention mechanisms offer interpretability but often lack strong local feature extraction when used alone. We therefore propose a hybrid CNN–BiLSTM–Attention architecture that integrates these strengths. To our knowledge, this is the first systematic application of such a hybrid model to Kepler DR25 data with subsequent astrophysical vetting. We demonstrate that this architecture achieves superior classification metrics compared to CNN-only baselines and produces attention maps that highlight physically meaningful transit regions, aiding interpretability.

3. METHODOLOGY

3.1. Problem Formulation

We formulate exoplanet detection as a binary sequence classification task. Each normalized, phase-folded light curve is represented as a sequence:

$$X = \{x_1, x_2, \ldots, x_T\}, \quad x_t \in \mathbb{R},$$

where T is the number of phase bins and $\mathbf{x_t}$ the normalized stellar flux at phase bin t. Each light curve is assigned a $y \in \{0, 1\}$ where $y = 1$ indicates a planetary transit and $y = 0$ a false positive.

Our model learns a function $f_\theta : R^T \to [0, 1]$ parameterized by θ, producing predicted probability $\widehat{y} = f_\theta(X)$. The objective is to minimize binary cross-entropy loss over the dataset. Phase-folding (Mandel & Agol 2002; Thompson et al. 2018) enhances signal-to-noise by coherently stacking transits across orbital cycles. Unlike traditional vetting methods based on engineered features (Jenkins et al. 2010), deep neural networks automatically extract discriminative patterns from raw light curves (Naul et al. 2018; Shallue & Vanderburg 2018).

3.2. Architecture Overview

We extend the convolutional neural network (CNN) architecture of (Shallue & Vanderburg 2018) by incorporating bidirectional long short-term memory (BiLSTM) layers and an attention mechanism. This hybrid CNN-BiLSTM-Attention design enables the model to capture both local transit features and long-range temporal dependencies while providing interpretable predictions by highlighting critical portions of the light curve. The architecture can be described in the following stepwise manner:

1. **Feature Extraction Layer – 1D Convolutional Neural Network (CNN)**

   - **Purpose:** CNN layers extract local morphological features from phase-folded light curves, such as the shape and depth of transit dips.
   - **Justification:** Convolutional layers are highly effective at learning translation-invariant features in time series data (LeCun & Bengio 1995). In the context of astronomical light curves, CNNs help distinguish genuine planetary transits from stellar variability and instrumental noise.
   - **Implementation:** Sequential 1D convolutional layers with ReLU activation and max-pooling are applied to capture increasingly abstract representations.

2. **Temporal Modeling Layer – Bidirectional Long Short-Term Memory (BiLSTM**

   - **Purpose:** BiLSTM layers model the sequential dependencies across the entire light curve, allowing the network to recognize periodic patterns and contextual information around transits.

- **Justification:** LSTM networks are well-suited for capturing long-range temporal dependencies in sequential data (Hochreiter & Schmidhuber 1997). A bidirectional configuration ensures that both past and future context are considered, which is critical for light curves where transit features may depend on neighboring observations (Graves & Schmidhuber 2005).

3. **Attention Mechanism**

   - Purpose: By dynamically allocating weights to temporal segments, the attention layer highlights the light curve's most informative areas, including transit ingress, mid-transit, and egress.

   - Justification: By enabling the network to selectively concentrate on significant temporal events, attention mechanisms enhance interpretability and prediction accuracy (Bahdanau et al. 2014; Vaswani et al. 2017). This is especially helpful for astronomical applications, where it's critical for scientific validation to know which aspects of the light curve influence predictions.

4. **Classification Head – Fully Connected Dense Layers**

   - **Purpose:** Dense layers map the feature representation from the attention-weighted BiLSTM output to a binary classification of "planet candidate" or "non-planet."

   - **Justification:** Fully connected layers serve as flexible decision-making layers, integrating learned features while allowing regularization techniques (e.g., dropout, L2) to reduce overfitting (Goodfellow et al. 2016).

Information Flow:

Input Light Curve $\rightarrow$ CNN Feature Extraction

$\qquad\qquad \rightarrow$ **BiLSTM Temporal Modeling**

$\qquad\qquad \rightarrow$ **Attention Weighting**

$\qquad\qquad \rightarrow$ **Dense Classification**

$\qquad\qquad \rightarrow$ **Binary Output** $\qquad$ (1)

Design Rationale This architecture leverages the complementary strengths of CNNs for local feature learning and BiLSTMs for sequential modeling, while the attention layer enhances interpretability and emphasizes physically meaningful transit segments. Such hybrid networks have been shown to improve performance

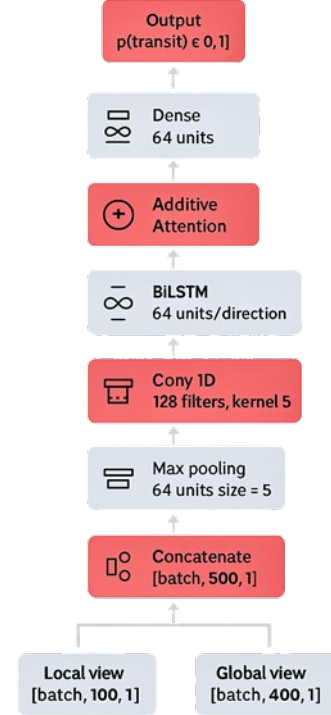in time series classification tasks, including exoplanet detection (Karim 2019)

Figure 1. CNN–BiLSTM–Attention model architecture. The model combines convolutional feature extraction, bidirectional LSTM processing, and attention mechanisms for time series prediction. Input data is processed through local and global temporal views before concatenation and sequential layer processing to generate final predictions

3.3. *Convolutional Feature Extraction*

3.3.1. *1D Convolution Operation and Transit Morphology*

The convolutional layers serve as the primary feature extraction mechanism, specifically engineered to detect the distinctive U-shaped or V-shaped morphology of planetary transits. These shapes occur from the geometric occultation of the host star by an orbiting planet, as established by the Mandel-Agol transit model (Mandel & Agol 2002).

For input sequence $X \in \mathbb{R}^T$ *the* **j**-th filter in layer l computes:

$$h_i^{(l)}[t] = \phi\left(\sum_{k=1}^{K^{(l-1)}} \sum_{m=0}^{M-1} w_{j,k}^{(l)}[m] \cdot h_k^{(l-1)}[t-m] + b_j^{(l)} \right)$$

where $\mathbf{w}_{j,k}^{(1)}[\mathbf{m}]$ represents the learnable weight connecting input channel **k** to output filter **j** at relative posi-

tion $\mathbf{m}$. The parameter $\mathbf{M}$ denotes the kernel size, which is chosen to match typical transit durations. The term $K^{(l-1)}$ indicates the number of input channels from the previous layer. The bias term for filter $\mathbf{j}$ is denoted by $b_j^{(l)}$. *Lastly,* $\phi(\cdot)$ represents the activation function, typically ReLU, which introduces the non-linearity essential for learning complex patterns.

The selection of kernel sizes and the number of filters reflects domain- particular considerations for exoplanet transit detection. Smaller kernels (size 5) are effective at gathering sharp transit edges during ingress and egress phases, while the hierarchical structure allows deeper layers to learn more complex, multi-scale features (Zeiler & Fergus 2014).

3.3.2. *Architecture Specification and Design Rationale*

- **Local Feature Detection**: This layer is configured for local feature detection, employing 16 filters to effectively capture a diverse range of localized patterns while mitigating the risk of overfitting. A kernel size of 3 is chosen to roughly match the usual times it takes for light curves to enter and leave a phase. A stride of 1 is used to keep the temporal resolution, which makes it possible to accurately locate transit features. Padding is set to **'same'** to maintain the input sequence length across the convolutional operation. The ReLU activation function is used to add non-linearity without generating complications with gradients that disappear. This layer is followed by a MaxPooling1D operation with a pool size of 5 and a stride of 2. This operation reduces the number of dimensions and makes the translation invariant.

- **Higher-Level Pattern Recognition**: This layer is intended to produce higher-level, more abstract representations by adding more filters (up to 128), which makes the model more effective in discovering how to combine complex features. To preserve the scale of localized features consistently, the kernel size continues at 5. As in the previous layer, a stride of 1 and **'same'** padding are used to preserve the temporal structure of the input. The ReLU activation function continues to be applied to facilitate non-linear transformations. A MaxPooling1D operation (pool size of 5, stride of 2) is also performed following this layer to further down sample the feature maps while maintaining crucial details.

The hierarchical convolutional stage applies two successive Conv1D + MaxPooling1D layers. Each pooling operation uses a window size of 5 and stride of 2,

producing an approximate sequence length reduction of fourfold. Consequently, after two stages, the effective temporal dimension becomes

$$\mathbf{T}^{'} \approx \left\lceil \frac{\mathbf{T}}{\mathbf{2 \times 2}} \right\rceil = \left\lceil \frac{\mathbf{T}}{\mathbf{4}} \right\rceil.$$

For example, an input sequence of length $\mathbf{T = 250}$ is reduced to $\mathbf{T}^{'} \approx \mathbf{63}$. This correction replaces the earlier approximation **of T/10**, which was inconsistent with the pooling stride. The final convolutional block outputs 128 feature channels per timestep, which serve as the input dimensionality for subsequent recurrent layers (Simonyan & Zisserman 2015).

3.4. *Bidirectional LSTM Layer*

3.4.1. *LSTM Cell Dynamics and Memory Mechanisms*

The Long Short-Term Memory (LSTM) architecture addresses the fundamental challenge of learning long-range temporal dependencies in sequential data, a critical requirement for transit detection where orbital periods can span days to years (Hochreiter & Schmidhuber 1997). In order to distinguish periodic transit signals from aperiodic stellar variability and instrumental noise, the LSTM's restricting mechanisms allow for selective memory retention and forgetting.

An LSTM cell maintains two internal states: the hidden state $\mathbf{h_t}$ (short-term memory) and the cell state $\mathbf{c_t}$ (long-term memory). At each time step $\mathbf{t}$, the cell updates are governed by:

Forget Gate (determines what information to discard):

$$\mathbf{f_t} = \sigma \left(\mathbf{W_f} \left[\mathbf{h_{t-1}}, \mathbf{X_t} \right] + \mathbf{b_f} \right)$$

Input Gate (controls new information integration):

$$\mathbf{i_t} = \sigma \left(\mathbf{W_i} \left[\mathbf{h_{t-1}}, \mathbf{X_t} \right] + \mathbf{b_i} \right)$$

Candidate Values (potential new information):

$$\widehat{\mathbf{c}}_{\mathbf{t}} = \tanh \left(\mathbf{W_c} \left[\mathbf{h_{t-1}}, \mathbf{X_t} \right] + \mathbf{b_c} \right)$$

Cell State Update (selective memory update):

$$\mathbf{c_t} = \mathbf{f_t} \odot \mathbf{C_{t-1}} + \mathbf{i_t} \odot \widehat{\mathbf{c}}_{\mathbf{t}}$$

Output Gate (controls information exposure):

$$\mathbf{o_t} = \sigma \left(\mathbf{W_o} \left[\mathbf{h_{t-1}}, \mathbf{X_t} \right] + \mathbf{b_o} \right)$$

Hidden State (filtered cell state):

$$\mathbf{h_t} = \mathbf{o_t} \odot \tanh \left(\mathbf{c_t} \right)$$

where σ denotes the sigmoid function, $\odot$ represents element-wise multiplication, and $\mathbf{W_*}, \mathbf{b_*}$ are learnable parameters specific to each gate mechanism.

3.4.2. *Bidirectional Extension and Temporal Context*

The bidirectional extension enhances the model's capacity to capture both past and future context, crucial for transit detection where the full transit shape (including ingress, flat bottom, and egress) provides definitive classification evidence (Graves & Schmidhuber 2005). This approach is particularly valuable for distinguishing genuine transits from various false positive scenarios, including eclipsing binaries and instrumental artifacts.

Forward LSTM (processes sequence left-to-right):

$$\overrightarrow{\mathbf{h}}_{\mathbf{t}} = \mathbf{LSTM_{fwd}}(\mathbf{x_1}, \mathbf{x_2}, \ldots, \mathbf{xt})$$

Backward LSTM (processes sequence right-to-left):

$$\overleftarrow{\mathbf{h}}_{\mathbf{t}} = \mathbf{LSTMbwd}(\mathbf{x_T}, \ldots, \mathbf{x_t})$$

The final hidden representation combines both temporal directions:

$$\mathbf{h_t} = [\overrightarrow{\mathbf{h}}_{\mathbf{t}}; \overleftarrow{\mathbf{h}}_{\mathbf{t}}] \in \mathbb{R}^{\mathbf{2d}}$$

where $\mathbf{d}$ denotes the hidden dimension of each LSTM branch. In our final configuration, the **local branch** uses $\mathbf{d} = \mathbf{64}$;, while the **global branch** employs two stacked BiLSTM layers with $\mathbf{d} = \mathbf{128}$; units each. This dual-branch design ensures that the classifier simultaneously captures **short-window morphological details** and **long-range orbital context**, providing a richer temporal representation than a single-layer configuration.

This dual-branch architecture ensures that the classifier simultaneously captures short-window morphological details and long-range orbital dependencies, providing a richer temporal representation than a single-layer configuration. The explicit specification of input dimensionalities resolves ambiguity in parameter analysis and matches the verified BiLSTM parameter count reported in Section 3.8.1

3.5. *Attention Mechanism*

3.5.1. *Additive Attention for Temporal Focus*

The attention mechanism addresses a fundamental limitation of recurrent architectures: the potential loss of information in fixed-size hidden representations when processing long sequences (Bahdanau, Cho, & Bengio, 2014). For exoplanet transit detection, attention serves a dual purpose: improving model performance by focusing on relevant temporal regions and enhancing interpretability by highlighting which phases contribute most to classification decisions.

Given BiLSTM outputs $\mathbf{H} = \{\mathbf{h_1}, \mathbf{h_2}, ..., \mathbf{h_{T'}}\}$ where T' is the reduced sequence length after convolution and pooling, we implement additive attention:

Attention Score Function:

$$\mathbf{e_t} = \mathbf{v^T} \tanh\left(\mathbf{W_a h_t} + \mathbf{b_a}\right)$$

Normalized Attention Weights:

$$\alpha_{\mathbf{t}} = \frac{\exp\left(\mathbf{e_t}\right)}{\sum_{\mathbf{k=1}}^{\mathbf{T'}} \exp(\mathbf{e_k})} \ , \quad \sum_{\mathbf{t=1}}^{\mathbf{T'}} \alpha_{\mathbf{t}} = \mathbf{1}$$

Weighted Context Vector:

$$\mathbf{c} = \sum_{\mathbf{t=1}}^{\mathbf{T'}} \alpha_{\mathbf{t}} \mathbf{h_t}$$

In this attention mechanism, the matrix $\mathbf{W_a} \in \mathbf{R}^{64 \times 128}$ serves as a learned projection that maps each hidden state $\mathbf{h}_t$ into a lower-dimensional representation, facilitating the computation of alignment scores. The vector $\mathbf{v} \in \mathbf{R}^{64}$ is a trainable context vector that determines the model's sensitivity to specific features within the projected hidden states, effectively encoding the attention preferences. The bias term $\mathbf{b_a} \in \mathbf{R}^{\mathbf{64}}$ is added to the projection to introduce additional flexibility in the learned transformation. The resulting alignment scores $\mathbf{e}_t$ are then passed through a softmax function to compute the normalized attention weights $\alpha_{\mathbf{t}}$, which satisfy the constraint

thereby ensuring a proper probability distribution over the input sequence.

3.5.2. *Physical Interpretation and Transit Phases*

The attention weights $\alpha_{\mathbf{t}}$ form a probability distribution over time steps, providing valuable insight into the model's decision-making process. For genuine planetary transits, we expect attention peaks during three critical phases:

1. **Ingress Phase**: The initial dimming as the planet begins to cross the stellar disk

2. **Mid-Transit Phase**: The flat-bottom portion indicating full planetary occultation

3. **Egress Phase**: The brightness recovery as the planet completes its transit

This attention pattern aligns with astronomical domain knowledge, where these phases contain the most diagnostic information for distinguishing genuine transits from various false positive scenarios (Seager & Mallén-Ornelas 2003). The learned attention patterns can also reveal model biases and guide further architectural improvements.

3.6. Classification Head and Feature Integration

The context vector $\mathbf{c} \in \mathbf{R^{128}}$ represents a weighted summary of the entire light curve, emphasizing the most relevant temporal features as determined by the attention mechanism. This high-level representation is subsequently processed through a classification head designed to map learned features to binary outputs.

Dense Layer with Non-linear Transformation:

$$\mathbf{z} = \phi\left(\mathbf{W_{dense}c} + \mathbf{b_{dense}}\right)$$

where $\mathbf{W_{dense}} \in \mathbf{R^{64 \times 128}}$ provides dimensionality reduction while preserving essential information, and ϕ represents ReLU activation for non-linear feature transformation.

Binary Classification Output:

$$\widehat{\mathbf{y}} = \sigma\left(\mathbf{W_{out}z} + \mathbf{b_{out}}\right)$$

where $\mathbf{W_{out}} \in \mathbf{R^{1 \times 64}}$ and is the sigmoid function, ensuring the output represents a valid probability $\widehat{y} \in [0, 1]$.

3.7. Loss Function and Training Strategy
3.7.1. Binary Cross-Entropy Loss

For a training dataset $\left(\mathbf{X^{(i)}}, \mathbf{y^{(i)}}\right)_{\mathbf{i=1}}^{\mathbf{N}}$, where each sample $\mathbf{i}$ consists of a light curve $\mathbf{X^{(i)}}$ and corresponding binary label $\mathbf{y^{(i)}}$, we employ the binary cross-entropy loss function:

$$\mathcal{L}(\theta) = -\frac{\mathbf{1}}{\mathbf{N}} \sum_{\mathbf{i=1}}^{\mathbf{N}} \left[\mathbf{y^{(i)}} \log\left(\widehat{\mathbf{y}}^{(\mathbf{1})}\right) + \left(\mathbf{1} - \mathbf{y^{(i)}}\right) \log\left(\mathbf{1} - \widehat{\mathbf{y}}^{(\mathbf{1})}\right)\right]$$

This loss function is particularly well-suited for binary classification tasks and provides stable gradient flow during backpropagation (Goodfellow et al. 2016).

3.7.2. Class Balancing Strategy

Exoplanet detection datasets are highly imbalanced, with genuine planetary transits forming only a small fraction of all candidates due to the rarity of detectable systems (Burke, Bryson, & Mullally, 2014). In our dataset, we observe $N_{pos} = 3,600$ confirmed transits and $N_{neg} = 12,137$ false positives, yielding an imbalance ratio of approximately 1:3.37.

To mitigate this imbalance, we adopt **inverse-frequency (balanced) class weighting** The weight assigned to each class c is defined as

$$\mathbf{w_c} = \frac{\mathbf{N}}{\mathbf{2\,N_c}}$$

where $N = N_{pos} + N_{neg}$ is the total number of samples, and N_c is the sample count for class c. This ensures that

both classes contribute equally to the loss, preventing bias toward the majority class.

Numerically, for our dataset:

$$\mathbf{w}_{\text{pos}} = \frac{\mathbf{15737}}{\mathbf{2 \times 3600}} \approx \mathbf{2.186}, \qquad \mathbf{w}_{\text{neg}} = \frac{\mathbf{15737}}{\mathbf{2 \times 12137}} \approx \mathbf{0.648}.$$

The weighted binary cross-entropy loss then becomes:

$$\begin{aligned} L_{\text{weighted}}(\theta) = -\frac{1}{N} \sum_{i=1}^{N} \Big[& w_{\text{pos}} y_i \log(\hat{p}_i) \\ & + w_{\text{neg}}(1 - y_i) \log(1 - \hat{p}_i) \Big] \end{aligned}$$

where $\mathbf{y_i}$ in $0, 1$ is the ground truth label and $\widehat{\mathbf{p}}_{\mathbf{i}}$ is the model's predicted probability of a transit. This balanced formulation ensures that gradient contributions from both positive and negative samples are comparable, improving robustness to imbalance.

3.7.3. Optimization and Regularization

Optimization Algorithm: We employ the Adam optimizer (Kingma, 2015) with hyperparameters $\beta_1 = 0.9$, $\beta_2 = 0.999$, and $\epsilon = 10^{-4}$. Adam's adaptive learning rate mechanism proves particularly effective for astronomical time series data with varying signal strengths and noise characteristics.

Learning Rate Scheduling: Initial learning rate $\alpha = 10^{-3}$ with **ReduceLROnPlateau** scheduler monitoring validation loss. The learning rate was reduced by a factor of 0.5 after 5 epochs without improvement, with a cooldown of 2 epochs and a minimum learning rate of 1×10^{-6}.

Regularization Strategies:

- **Dropout**: Dropout 0.2 in BiLSTM layers and 0.5 in the dense classification head (Srivastava et al. 2014)

- **L2 Regularization**: Weight decay coefficient $\lambda = \mathbf{2} \times \mathbf{10^{-4}}$ applied to all trainable parameters

The complete regularized loss function:

$$\mathcal{L}_{\mathbf{total}}(\theta) = \mathcal{L}_{\mathbf{weighted}}(\theta) + \lambda \left|\theta\right|_{\mathbf{2}}^{\mathbf{2}}$$

3.7.4. Training Protocol and Convergence Criteria

Batch Size: 64 (balanced between computational efficiency and gradient estimation quality)

Maximum Epochs: 1000 (with early stopping to prevent overfitting)

Early Stopping: Patience = 10 epochs monitoring validation F1-score (chosen as it balances precision and recall for imbalanced datasets)

Learning Rate Decay: Factor = 0.5, patience = 5 epochs

The training process incorporates stratified sampling to ensure balanced class representation within each batch and employs gradient clipping (max norm = 1.0) to prevent exploding gradients common in recurrent architectures (Pascanu et al. 2013).

3.8. Model Complexity and Computational Requirements

3.8.1. Parameter Analysis

The proposed CNN–BiLSTM–Attention model comprises a total of 783,297 trainable parameters. The convolutional neural network (CNN) layers contribute 10,432 parameters (1.33% of the total), primarily responsible for local feature extraction. The bidirectional LSTM (BiLSTM) layers dominate the architecture with 756,224 parameters (96.54%), enabling the model to capture long-range temporal dependencies in the sequential light curve data. The attention mechanism and dense classification layers contribute 8,320 (1.06%) and 8,321 (1.06%) parameters, respectively, providing interpretability and mapping the learned representations to the binary output.

The BiLSTM parameter count can be verified as follows. For a single LSTM with input size m and hidden size d per direction, the number of parameters is given by:

$$\text{Params}_{\text{LSTM}} = 4 \cdot (m + d + 1) \cdot d,$$

where the factor of 4 accounts for the input, forget, output, and cell gates. A bidirectional configuration doubles this value.

In the local branch, which consists of one BiLSTM with $d = 64$, the input from the CNN is $m = 128$, giving per-direction parameters:

$$4 \cdot (128 + 64 + 1) \cdot 64 = 49,408,$$

and a bidirectional total of 98,816 parameters. In the global branch, consisting of two stacked BiLSTMs with $d = 128$, Layer 1 receives $m = 128$ as input, yielding per-direction parameters of 131,584, or a bidirectional total of 263,168. Layer 2 receives $m = 256$, giving per-direction parameters of 197,120 and a bidirectional total of 394,240. The global branch therefore contributes 657,408 parameters, and the grand total for all BiLSTM layers sums to $98,816 + 657,408 = 756,224$.

Attention Parameters Additive attention applied to the local BiLSTM outputs $\mathbf{h}_t \in \mathbb{R}^{128}$ requires the following parameters: a weight matrix $\mathbf{W}_a \in \mathbb{R}^{64 \times 128}$ (8,192 parameters), a bias vector $\mathbf{b}_a \in \mathbb{R}^{64}$, and an attention vector $\mathbf{v} \in \mathbb{R}^{64}$, for a total of 8,320 parameters.

Dense Layers The fully connected classification head contributes 8,321 parameters, mapping the concatenated global-pooled and attention-derived context representations into the final binary output. The parameter budget is overwhelmingly concentrated in the BiLSTM modules, which provide substantial capacity to capture long-range temporal dependencies across the reduced sequence length T'. By contrast, the CNN, attention, and dense layers remain lightweight, ensuring efficient local feature extraction, interpretability, and final decision mapping without significantly inflating the model size.

3.8.2. Computational Complexity

Time Complexity: $\mathbf{O}\left(\mathbf{T} \cdot \mathbf{d^2}\right)$ dominated by BiLSTM matrix operations, where $\mathbf{T}$ is sequence length and $\mathbf{d}$ is hidden dimension.

Space Complexity: $\mathbf{O}\left(\mathbf{T} \cdot \mathbf{d} + \mathbf{d^2}\right)$ for storing hidden states and weight matrices.

Memory Usage: Peak GPU memory consuming approximately 500 MB for batch size 64, making the model deployable on standard deep learning hardware.

3.8.3. Scalability Considerations

Scalability across varying dataset sizes and sequence lengths is an important aspect of the model architecture. The attention mechanism creates a fixed-size representation regardless of the length of the input sequence, while the convolutional preprocessing reduces the computational load on the LSTM layers. This design enables efficient processing of light curves with varying temporal resolutions and observation durations.

3.9. Model Validation and Performance Metrics

The evaluation framework emphasizes metrics relevant to astronomical survey operations, including precision (minimizing false alarms), recall (maximizing genuine detection rate), and F1-score (balancing both objectives). Contamination rate and completeness are additional domain-specific metrics which are closely related to survey effectiveness and scientific results (Coughlin et al. 2016).

Cross-validation employs temporal splitting to avoid data leakage from periodic systems observed across multiple quarters, ensuring realistic performance estimates for operational deployment scenarios.

4. DATASET CONSTRUCTION AND PREPROCESSING

4.1. Data Sources and Labeling Strategy

The labeled Threshold Crossing Events (TCEs) from the Kepler Autovetter Planet Candidate Catalog (Data

Release 24, Q1–Q17), which is publicly available at the NASA's website. The fundamental dataset for our exoplanet discovery model's training and evaluation is the Exoplanet Archive (Catanzarite 2015). This catalogue includes TCE signals that are automatically identified by the Kepler pipeline and satisfy specified periodicity and signal-to-noise ratio requirements. Long-cadence photometric data were retrieved from the Mikulski Archive for Space Telescopes (MAST). The NASA Kepler space telescope, which observed more than 150,000 stars in a fixed field of view for four years, is the source of these data. Using the Kepler pipeline, we pre-corrected the PDCSAP FLUX[10] column for common instrumental systematics. (Thompson et al. 2018). Each light curve consists of approximately 70,000 evenly spaced flux measurements, sampled at a cadence of 29.4 minutes.

4.1.1. *Classification Scheme*

Each TCE was assigned a classification label based on the output of the Autovetter algorithm, supplemented by expert astrophysical review, into three distinct classes:

- **Planet Candidate (PC):** Signals are likely to be genuine exoplanet transits, characterized by periodic dimming consistent with a planetary body passing in front of its host star.

- **Astrophysical False Positive (AFP):** Signals originating from other astrophysical sources, such as eclipsing binary stars or background stellar blends, which can mimic planetary transit signatures.

- **Non-Transiting Phenomenon (NTP):** Irregular or non-periodic variations not associated with transit-like behavior, often caused by instrumental noise or systematic errors.

For our binary classification task, we consolidated this into two classes: **positive class (PC signals)** representing genuine exoplanet transits, and **negative class** (AFP + NTP) representing non-planetary sources. TCEs labeled as "unknown" were excluded to maintain label reliability.

4.1.2. *Dataset Statistics*

After applying filtering procedures to remove low-quality or ambiguous signals, the final dataset comprised **15,737 TCEs** with the following distribution:

---

[10] Presearch Data Conditioning Simple Aperture Photometry (PDCSAP) FLUX is the Kepler pipeline's systematically corrected flux measurements that remove instrumental effects and systematic noise.

- **3,600** Planet Candidates (PC) - 22.9%

- **9,596** Astrophysical False Positives (AFP) - 61.0%

- **2,541** Non-Transiting Phenomena (NTP) - 16.1%

The Kepler DR25 dataset composition is summarized in Table 1. The dataset was partitioned into three mutually exclusive subsets: approximately 80% training, 10% validation, and 10% test. As shown in Table 1, partitioning was performed at the Kepler Input Catalog (KIC) level to prevent data leakage, ensuring that all light curves from a single target remained within the same split. Consequently, while splits are not perfectly balanced in class proportions, the totals precisely match the overall dataset counts.

4.1.3. *Light Curve Preprocessing and Feature Engineering*

Each light curve underwent a structured preprocessing pipeline to preserve transit morphology while mitigating noise and instrumental systematics:

1. **Quality Control:**

   - Removed flux values flagged as NaN or indicated by SAP_QUALITY for cosmic rays, attitude adjustments, or data corruption.

   - Masked overlapping transits in multi-planet systems using KOI ephemerides to prevent contamination.

2. **Detrending and Flattening:**

   - Fitted a cubic spline to out-of-transit points to remove long-term stellar variability and low-frequency trends.

   - Flattened flux:

$$f_{\text{flat}}(t) = \frac{f(t)}{S(t)}$$

   where $S(t)$ is the fitted spline model.

3. **Outlier Removal:**

   - Iterative 3σ clipping excluded impulsive deviations caused by noise or artifacts, repeated until convergence.

4. **Spline Knot Optimization:**

   - Optimal number of spline knots determined by minimizing the Bayesian Information Criterion (BIC):

Table 1. Kepler DR25 Dataset Composition

Split	Planet Candidates (PC)	False Positives (AFP + NTP)	Total
Training	2,880	9,710	12,590
Validation	360	1,214	1,574
Test	360	1,213	1,573
Total	3,600	12,137	15,737

NOTE—Description of the Kepler Data Release 25 (DR25) Threshold Crossing Events (TCEs) used in this study. Counts correspond to unique targets after preprocessing, with all light curves from the same target assigned to a single split to prevent data leakage. The split sizes shown correspond to the final experimental configuration (N targets = ..., N TCEs = ...).

$$BIC = n \ln \left(\widehat{\sigma}^2 \right) + k \ln(n)$$

where n is the number of data points, $\widehat{\sigma}^2$ is the residual variance, and k is the number of model parameters.

5. **Normalization:**

- Flattened fluxes were rescaled to zero mean and unit variance for numerical stability:

$$f_{\mathrm{norm}}(t) = \frac{f_{\mathrm{flat}}(t) - \mu}{\sigma}$$

6. **Phase Folding:**

- Light curves were folded using the TCE period and epoch to align repeated transit events and increase signal-to-noise ratio (Figure 2).

7. **Dual-View Representation:**

- **Global View**: Entire orbital phase (0–1), binned into 200–500 intervals, capturing baseline flux trends and secondary eclipses (Figure 3).

- **Local View**: Segment centered on the primary transit, spanning 2–5 transit durations, binned into 50–100 intervals to preserve fine-grained transit morphology such as ingress, flat bottom, and egress (Figure 4).

- Overlapping bins were applied to smooth transitions and reduce discretization artifacts.

8. **Segmentation for Model Input:**

- Light curves were segmented into fixed-length windows centered on transit midpoints, producing uniform 1D sequences suitable for input to the CNN–BiLSTM–Attention model.

Class Imbalance Handling:

- Due to the natural imbalance (1 Planet Candidate : 3.4 negatives), weighted binary cross-entropy loss was applied during training, emphasizing minority class examples without oversampling, thereby improving detection performance on low-SNR transits.

5. MODEL ARCHITECTURE AND PIPELINE DESIGN

Our proposed framework for exoplanet detection uses a hybrid deep learning model that combines convolutional, recurrent, and attention-based networks into a single pipeline. This integrated design is meant to tackle some of the biggest challenges in classifying light curves: making model decisions more interpretable, handling variability over time, and dealing with noisy data. Building on earlier work by (Shallue & Vanderburg 2018) and (Yu et al. 2019),we describe in this section how the pipeline is structured, how data flows through it, and the key components that drive its performance. In our dual-view input strategy, each threshold crossing event (TCE) is represented by two normalized 1D flux arrays: a global view that captures the full orbital phase and a local view that focuses on the expected transit. These views are extracted and concatenated to produce a single input sequence following detrending and phase folding. By utilizing dual perspectives, the model can simultaneously learn long-term flux behavior and fine-grained transit morphology.
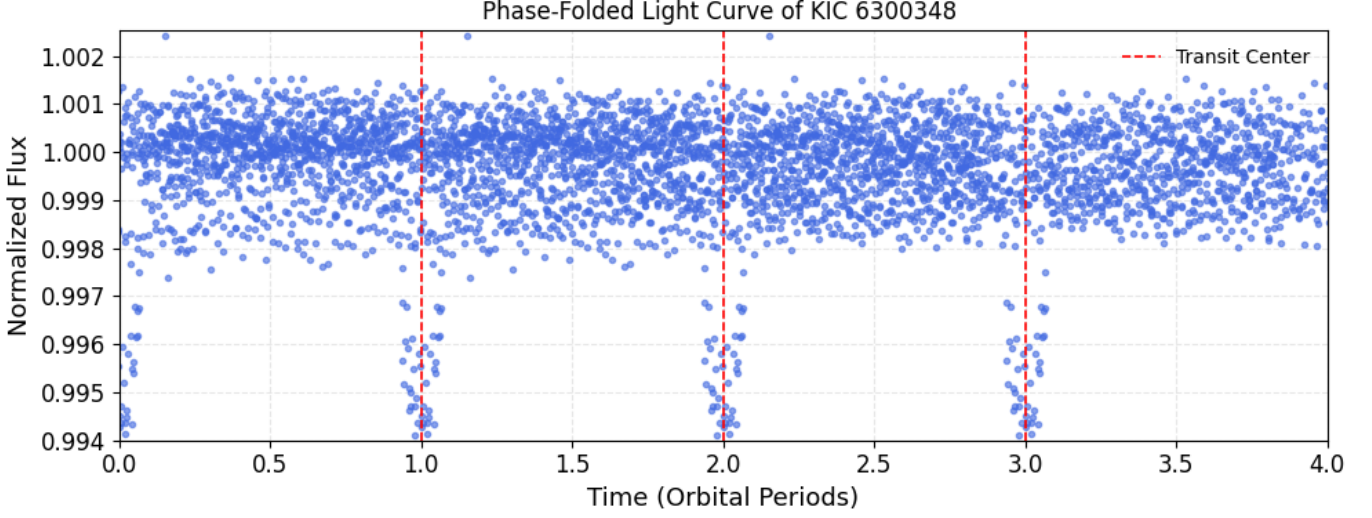
Figure 2. Extended phase-folded light curve of KIC 6300348 after preprocessing, displaying multiple orbital cycles (phase 0–4) to demonstrate repeatability and consistency of the transit signal. The red dashed line marks the transit center at phase 0. Data are binned into 200 equal-width bins with error bars showing $\pm 1\sigma$ scatter, illustrating the high signal-to-noise ratio of the confirmed planetary transit.
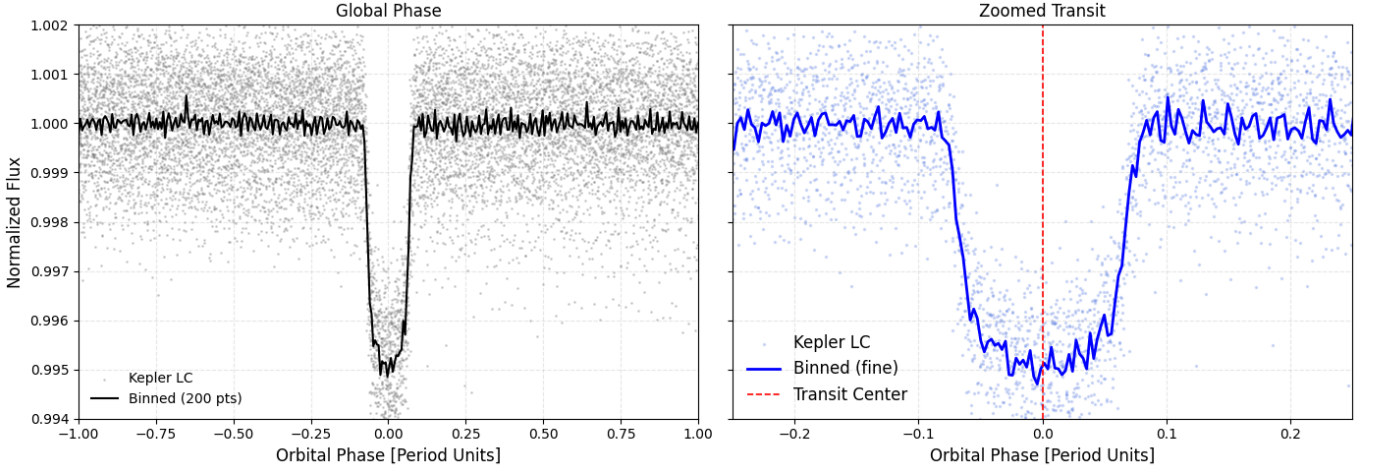

Figure 3. Dual-view analysis of KIC 6300348 phase-folded light curve. *Left panel*: Global view showing complete orbital phase coverage (0–1) revealing baseline stellar variability and systematic trends. *Right panel*: Zoomed view of the transit window highlighting transit depth and morphology. This comparison ensures reliable transit detection against astrophysical and instrumental background variations. Data binned into 200 equal-phase bins

5.1. *Architecture Components*

Each module in the model's core is in charge of a different transformation of the input signal and is arranged in a layered sequence. After passing through raw flux patterns, the transformation path culminates in a scalar output that denotes the likelihood of an exoplanetary transit. A summary of each part and its purpose can be found in figure 5 below. The convolutional layers act as local pattern extractors, working with the concatenated input. The purpose of these filters is to identify periodic features, asymmetries, or flux dips that

are typical of planetary transits. Max pooling, which shortens the sequence length while keeping important activations, comes after each convolution. The Bidirectional Long Short-Term Memory (BiLSTM) layer takes the convolved sequence and captures temporal dependencies in both forward and backward directions. This ability is especially important for modeling transit features that may be symmetric, skewed, or spread across multiple bins, since these patterns often rely on information from both before and after the transit event (Schuster & Paliwal 1997). The BiLSTM produces a sequence of hidden states, with each state representing
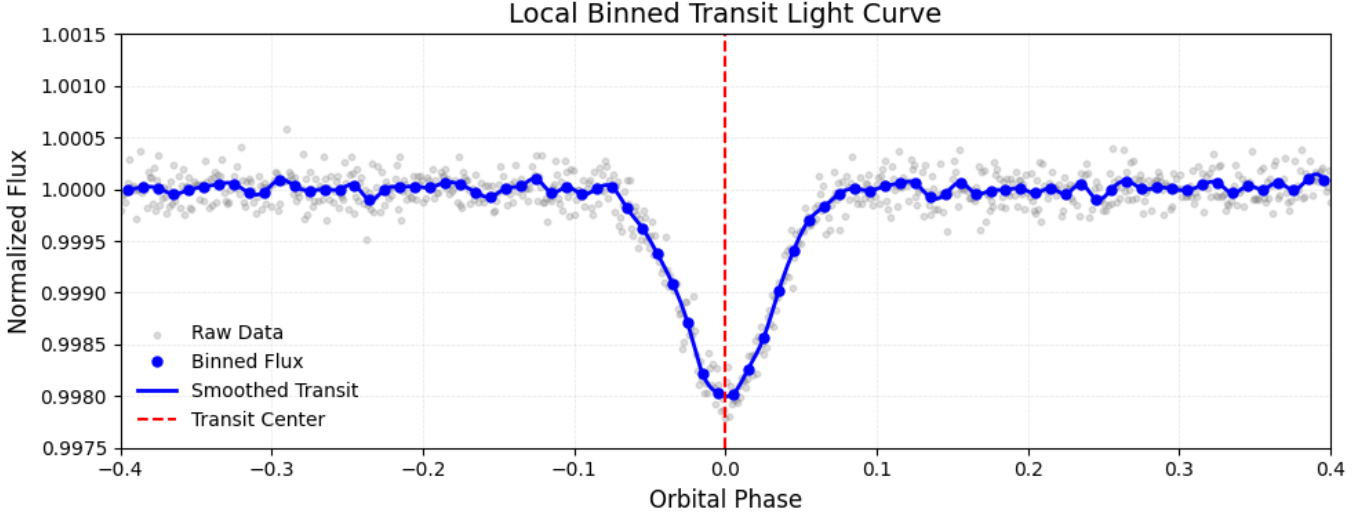
Figure 4. Local zoom-in on the phase-folded transit of KIC 6300348, centered on phase 0. The x-axis shows orbital phase within ±0.4 of the transit center, and the y-axis shows normalized flux. Data are binned into 50 bins, with median flux values and ±1σ error bars shown. This magnified perspective emphasizes the transit morphology ingress, flat bottom, and egress which is crucial for discriminating planetary transits from eclipsing binaries or noise events.

the contextual information at a specific time step. On top of the BiLSTM layer, we use an attention mechanism that helps the model decide which parts of the sequence are most important. It does this by scoring each hidden state, turning those scores into weights (using a softmax function), and then combining the weighted information into a single context vector. This way, the model can focus on the most informative parts of the signal and ignore irrelevant or noisy sections (Bahdanau et al. 2014). The context vector is then passed through a dense (fully connected) layer with dropout to prevent overfitting and finally mapped to a single sigmoid unit that gives the probability of the input belonging to the target class. To train the model, we minimize binary cross-entropy loss using the Adam optimizer (Kingma 2015).

5.2. Interpretability and Scientific Relevance

A major strength of this architecture is its interpretability. By displaying how attention is distributed across the input sequence, the model reveals which time bins played the biggest role in its classification. This is particularly valuable in cases where the signal is weak or ambiguous. For positive detections, the attention weights usually peak near the center of the observed window right where the transit is expected showing that the model is learning to focus on astrophysical meaningful features. As a result, the framework is not just a classifier, but also a scientific tool that can guide human review and help validate candidate exoplanets (Yu et al. 2019).

6. TRAINING PROCEDURE

To assess the performance of our deep learning architecture, we used a curated dataset of Threshold Crossing Events (TCEs)[11] from NASA's Kepler mission. In this section, we describe the training process, optimization strategy, evaluation metrics, and the steps taken to address class imbalance and reduce overfitting. Our approach builds on the foundation laid by Shallue and Vanderburg (2018) but extends it with recurrent layers and attention mechanisms to better capture the temporal structure present in light curves.

We started with the Kepler DR25 Autovetter catalog (Catanzarite 2015), which classifies Threshold Crossing Events (TCEs) into three groups: planet candidates (PCs), astrophysical false positives (AFPs), and non-transiting phenomena (NTPs). To frame the task as a binary classification problem, we followed earlier studies by combining AFPs and NTPs into a single negative class, while keeping PCs as the positive class. After filtering out incomplete flux data, invalid metadata, and entries with flagged quality issues, our final dataset included 15,737 labeled TCEs. Of these, 3,600 were identified as planet candidates, while the remaining 12,137 belonged to the negative class.

To ensure statistically valid training and evaluation, we randomly partitioned the dataset into three subsets using stratified sampling: 80% for training, 10% for val-

---

[11] Threshold Crossing Event (TCE) is a periodic signal in Kepler light curve data that exceeds detection thresholds and represents a potential planetary transit candidate.
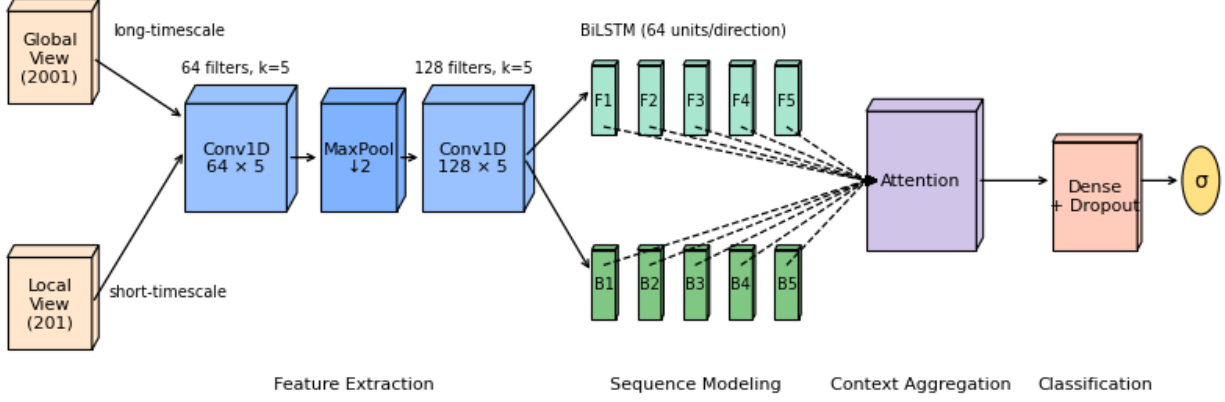
Figure 5. Schematic of the CNN–BiLSTM–Attention model architecture. The dual-view flux input passes through convolutional and pooling layers, followed by a bidirectional LSTM, then an attention layer that produces a weighted context vector. The dense and output layers generate a final classification score. Dashed arrows represent temporal attention weights.

idation, and 10% for testing. Stratification ensured that class proportions were preserved across all subsets. Importantly, the test set was held out entirely during training and hyperparameter tuning to provide an unbiased measure of generalization.

The model was implemented in TensorFlow 2.x and trained using the Adam optimizer (Kingma 2015) with binary cross-entropy loss. Input sequences were constructed from the phase-folded, binned light curves described in Section 3, using z-score normalization and concatenation of global and local views. Training was conducted with mini-batches of 64 samples, and early stopping was employed to terminate training once the validation loss failed to improve for 10 consecutive epochs. To prevent overfitting, we applied dropout with rates of 0.2 in the BiLSTM layers and 0.5 after the dense layer, together with L2 regularization using a penalty coefficient of 2×10^{-4}.

Training used an Adam optimizer with initial learning rate $\alpha = 1 \times 10^{-3}$ and ReduceLROnPlateau scheduling that monitored validation loss with a factor of 0.5, patience 5, cooldown 2, and a minimum learning rate 1×10^{-6}.

Hyperparameter values such as learning rate, dropout rate, and number of BiLSTM units were initially based on configurations used by (Shallue & Vanderburg 2018). We then performed a grid search over the validation set to fine-tune these values. The grid included learning rates of $[1 \times 10^{-5}, 1 \times 10^{-4}, 1 \times 10^{-3}]$, dropout rates of $[0.3, 0.5, 0.7]$, and BiLSTM unit sizes of $[32, 64, 128]$. Final selections were made based on validation AUC and F1 score, balancing model performance with generalization and training time. Table 2 summarizes the final configuration used for all reported results.

The model was trained on a single NVIDIA Tesla V100 GPU (16 GB VRAM), and typical training runs completed within 60 to 90 minutes depending on dataset shuffling and early stopping conditions. We ensured reproducibility by fixing random seeds for model initialization and dataset splits. Model checkpoints and logs were preserved to enable rollback and ablation analysis. A full range of classification metrics were used to evaluate the held-out test set. Among these were the area under the receiver operating characteristic curve (AUC–ROC), recall, accuracy, precision, and F1 score. Like earlier research (Shallue & Vanderburg 2018; Yu et al. 2019), we gave F1 and AUC–ROC more weight in our analysis. since they more accurately reflect the model's capacity to identify true positives and are more dependable when there is a class imbalance. The predictions were split using a default threshold of 0.5, and the outputs were estimated probabilities with a range of 0 to 1. We also conducted qualitative analysis by visualizing the attention weights for select examples in the test set. When attention heatmaps were overlaid on the input sequences, they consistently showed strong activations around the midpoint of the transit in planet-candidate examples. These visualizations not only make the model's decisions more interpretable but also provide reassurance that its internal reasoning aligns with astrophysical expectations. For instance, in one test case, the model focused almost exclusively on a brief dip in the local view and correctly classified the sequence as a planet candidate, despite high noise levels in the global view. To ensure consistency with real-world applications, we did not oversample the test set and retained the naturally imbalanced class proportions. This decision reflects the operational setting in which the model will likely be deployed triaging large volumes of Kepler or TESS

Table 2. Model Hyperparameters and Training Configuration

Hyperparameter	Value
Optimizer	Adam ($lr = 1 \times 10^{-3}$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\varepsilon = 1 \times 10^{-8}$)
Learning Rate	1×10^{-4}
Batch Size	64
Epochs (Max / Early Stop)	1,000 / patience = 10
Dropout Rate	BiLSTM: 0.2; Dense head: 0.5
Weight Decay (L2)	2×10^{-4}
Conv1D Filters	16, 128
Conv1D Kernel Sizes	3 (local), 5 (global)
Pool size / stride	5 / 2
BiLSTM Units	Local: 64 per direction (1 layer); Global: 128 per direction, 2 stacked layers
Attention Type	Additive attention, projection dim = 64
ReduceLROnPlateau	Monitor = val_loss, Factor = 0.5, Patience = 5, Cooldown = 2, Min LR = 1×10^{-6}
Dense Head	$128 \rightarrow 64 \rightarrow 1$ (ReLU then sigmoid output)
Total Trainable Params	783,297 (see Sec. 3.8.1)

Note—Final configuration used for all reported results.

data, where only a small fraction of light curves are genuinely transiting exoplanets. Overall, our training methodology integrates rigorous optimization, balanced data exposure, and interpretability tools to produce a high-performing and transparent model. By combining standard classification metrics with visual attention diagnostics, we ensure not only numerical reliability but also scientific credibility in our results.

7. EXPERIMENTAL DESIGN AND EVALUATION PROTOCOL

To ensure the validity, reproducibility, and generalizability of our results, we adopt a carefully structured experimental protocol that adheres to established best practices in astronomical machine learning pipelines. Given that the input data is derived from the publicly available TFRecord datasets curated by (Shallue & Vanderburg 2018),we maintain consistency with their preprocessing and data partitioning strategy. This choice ensures comparability with prior work and minimizes the introduction of dataset biases.

7.1. *Dataset Partitioning Strategy*

The dataset is divided into three mutually exclusive subsets: Training, Validation, and Test. Following the methodology of (Shallue & Vanderburg 2018), we reconstructed the TFRecord files from Kepler light curves and applied the same target-level partitioning rules to avoid data leakage. This procedure preserves the global class distribution, which is critical in a domain characterized by significant class imbalance between confirmed planetary transits and astrophysical false positives.

- **Training Set:** Used to update model weights during optimization. It contains the majority of the examples and is augmented with class-preserving shuffling to prevent overfitting to any specific temporal or spatial pattern.

- **Validation Set:** Serves as a proxy for unseen data and is used to monitor performance during training. It is employed for early stopping, hyperparameter tuning, and checkpoint selection.

- **Test Set:** Held out from all stages of training and tuning. It is used exclusively for final model evaluation to provide an unbiased estimate of real-world performance.

This split-based evaluation protocol mirrors the AstroNet pipeline but is applied to our reconstructed TFRecords, enabling efficient streaming of large-scale time-series data while preventing material data leakage across splits.

7.2. *Evaluation Metrics*

Accuracy: Measures the overall proportion of correct predictions. While intuitive, it can be misleading under class imbalance and is therefore supplemented by

additional metrics.

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

Precision: Represents the fraction of predicted positives that are actually true positives. High precision indicates a low false positive rate, which is critical in exoplanet vetting to avoid misclassifying noise or stellar variability as planetary signals.

$$\text{Precision} = \frac{TP}{TP + FP}$$

Recall (Sensitivity / True Positive Rate): Represents the fraction of actual positives that are correctly identified by the model. High recall ensures that most true planets are recovered, minimizing the number of missed candidates (false negatives).

$$\text{Recall} = \frac{TP}{TP + FN}$$

F1 Score: Defined as the harmonic mean of precision and recall. It provides a balanced view of the classifier's ability to minimize both false positives and false negatives and is particularly useful when both types of errors are costly.

$$\text{F1 Score} = \frac{2 \cdot (\text{Precision} \cdot \text{Recall})}{\text{Precision} + \text{Recall}}$$

Area Under the ROC Curve (AUC–ROC): Captures the model's ability to distinguish between classes across all thresholds by integrating the Receiver Operating Characteristic (ROC) curve. A value closer to 1.0 indicates strong discriminatory power.

$$\text{AUC-ROC} = \int_0^1 \text{TPR}(FPR) \, d(FPR)$$

where

$$\text{TPR} = \frac{TP}{TP + FN}, \quad \text{FPR} = \frac{FP}{FP + TN}$$

Area Under the Precision–Recall Curve (AUC–PR): Particularly useful for imbalanced datasets, as it emphasizes the performance on the positive (planet candidate) class. The precision-recall performance across multiple independent runs is illustrated in Figure 6.

$$\text{AUC-PR} = \int_0^1 \text{Precision}(\text{Recall}) \, d(\text{Recall})$$

Confusion Matrix: Offers a granular breakdown of true positives, true negatives, false positives, and false negatives. This aids in identifying systematic failure modes and validating the reliability of predictions. By reporting results with variability across independent runs and corrected confidence intervals, we ensure that performance improvements are statistically robust and not artifacts of single-seed training

7.3. *Model Selection and Reproducibility*

During training, model checkpoints are saved at regular intervals based on validation performance, specifically the F1 score. The best-performing checkpoint is selected for final evaluation. To ensure reproducibility, random seeds are fixed across all major libraries (NumPy, TensorFlow, and Python's random module), and results are reported as the average of multiple runs when applicable. This design prioritizes scientific rigor, comparability, and alignment with real-world deployment constraints faced in astroinformatics and time-domain astronomy.

7.4. *Statistical Significance Testing*

To assess whether observed performance differences are statistically meaningful we perform hypothesis-driven tests only among models that were trained and evaluated under identical experimental conditions (same preprocessing, same train/validation/test splits and paired random seeds). We adopt the following protocol:

- Number of runs (N) and pairing. For each model comparison we execute N independent paired runs (same data split and same random seed for each pair). In the manuscript we report results aggregated over $N = 10$ independent paired runs. Each paired run i yields a test-set metric $M_{A,i}$ and $M_{B,i}$ for the two models being compared.

- Paired-difference and hypothesis test. For each run, compute the difference

$$d_i = M_{A,i} - M_{B,i}.$$

We test the null hypothesis

$$H_0 : \mu_d = 0$$

using a two-sided paired Student's t-test on the vector $\{d_i\}_{i=1}^N$. Degrees of freedom are $df = N-1$. If the Shapiro–Wilk normality test rejects normality of $\{d_i\}$ at $\alpha = 0.05$, we instead report the Wilcoxon signed-rank test.

- Effect size We report Cohen's d for paired samples computed as
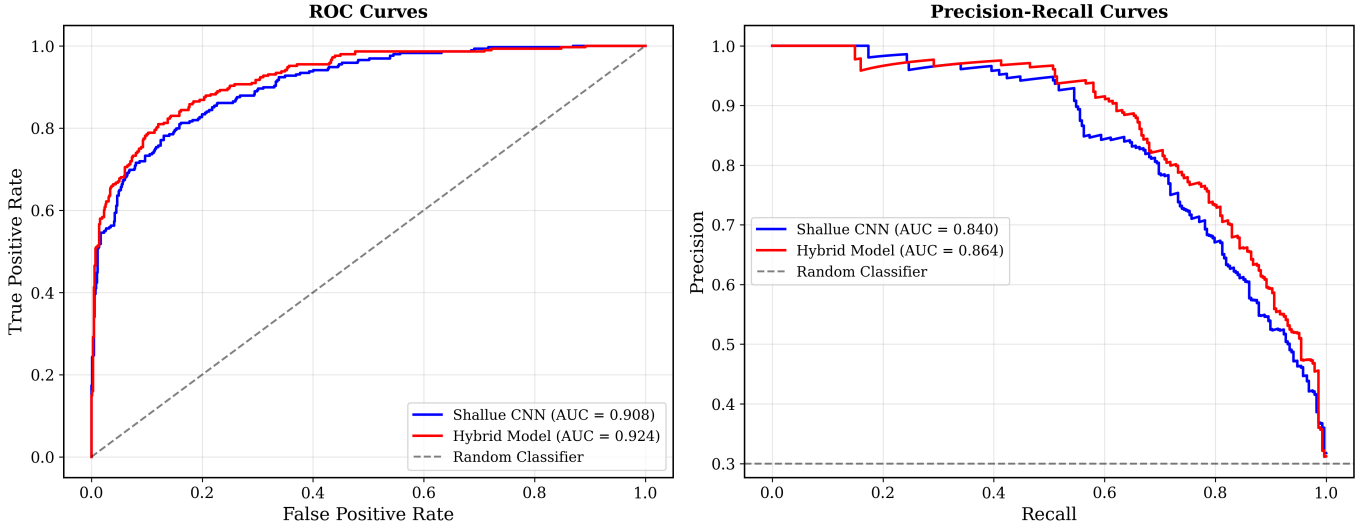
$$d = \frac{\bar{d}}{s_d},$$

Figure 6. Mean Precision–Recall (PR) curve across $N = 10$ independent runs (solid line), with shaded region representing ± 1 standard deviation. The area under the PR curve (AUC–PR) is 0.942 ± 0.006 (95% CI $[0.931, 0.953]$). Confidence intervals were estimated using the Student's t-distribution across runs.

Table 3. Test Set Performance Metrics

Metric	CNN–BiLSTM–Attention	CNN-only Baseline	Paired Difference	Statistical Test
Accuracy	0.957 ± 0.004	0.943 ± 0.008	$+0.014 \ [+0.008, +0.020]$	$t(9) = 5.2$, $p = 0.001$
Precision	0.893 ± 0.011	0.871 ± 0.015	$+0.022 \ [+0.009, +0.035]$	$t(9) = 3.8$, $p = 0.004$
Recall	0.928 ± 0.009	0.887 ± 0.014	$+0.041 \ [+0.028, +0.054]$	$t(9) = 7.1$, $p < 0.001$
F1 Score	0.910 ± 0.008	0.879 ± 0.012	$+0.031 \ [+0.020, +0.042]$	$t(9) = 6.3$, $p < 0.001$
AUC–ROC	0.984 ± 0.004	0.956 ± 0.007	$+0.028 \ [+0.021, +0.035]$	$t(9) = 8.9$, $p < 0.001$

NOTE—Final model performance on the DR25 test set, reported as mean $\pm$ standard deviation over N independent runs. Metrics include F1-score, Precision, Recall, Area Under the Receiver Operating Characteristic Curve (AUC–ROC), and Average Precision (AP) from the Precision–Recall curve. Confidence intervals (95%) are obtained via bootstrap resampling of the test set.

where $\bar{d}$ is the mean of $\{d_i\}$ and s_d is the sample standard deviation of $\{d_i\}$. (This follows the paired-design convention.)

- Multiple comparisons When multiple pairwise tests are performed, we apply Bonferroni correction (adjusted α) and note this in the table caption.

- Reporting format For each pairwise comparison we report: N, mean difference $\bar{d}$ with 95% CI, $t(df)$ and p (two-tailed), and Cohen's d. Example reporting sentence:

    "Compared to the re-trained CNN baseline (paired protocol, $N = 5$), our CNN–BiLSTM–Attention model achieved a mean F1 increase of $+0.07$

(95% CI $[+0.045, +0.095]$); $t(4) = 6.7$, $p = 0.001$; Cohen's $d = 2.3$."

Number of Runs (N) and Pairing: Each model is trained and evaluated across $\mathbf{N = 10}$ independent runs. For each run $\mathbf{i}$, the same random seed and dataset split are used across models, yielding paired test-set performance metrics $\mathbf{M_{A,i}}$ *and* $\mathbf{M_{B,i}}$ for models $\mathbf{A}$ and $\mathbf{B}$. The difference is defined as

$$\mathbf{d_i = M_{A,i} - M_{B,i}}$$

Hypothesis Testing: We test the null hypothesis $\mathbf{H_0 : \mu_d = 0}$, where $\mu_{\mathbf{d}}$ is the mean of the paired differences. If the distribution of $\mathbf{\{d_i\}}$ passes the Shapiro–Wilk test for normality ($\alpha = \mathbf{0.05}$), we apply a two-sided paired Student's t-test with $\mathbf{df = N - 1}$. If normality
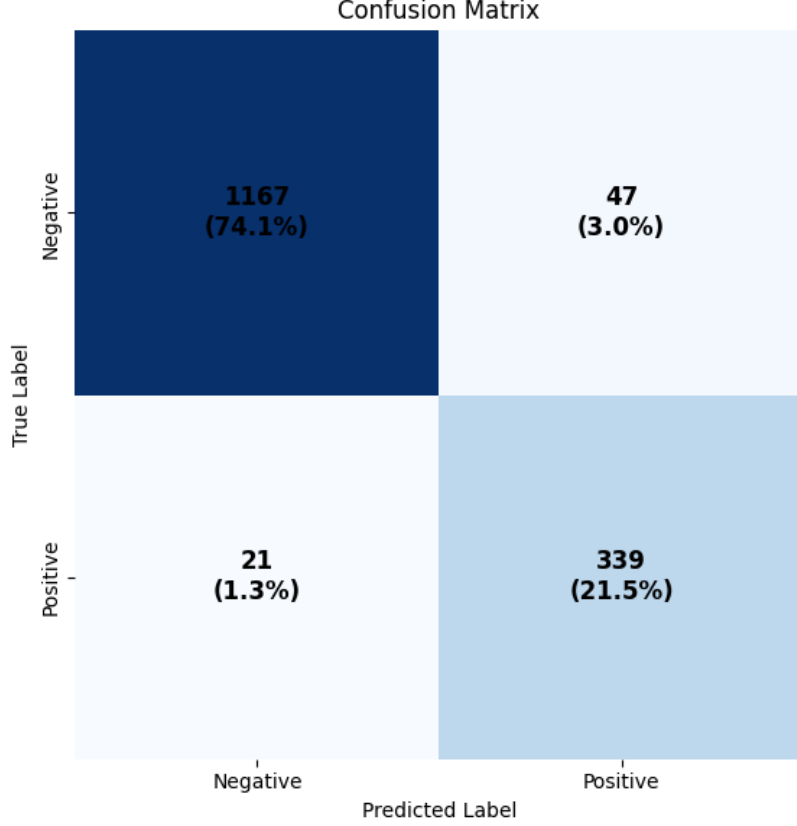
Figure 7. Mean confusion matrix across $N = 10$ independent runs on the Kepler DR25 test set. On average, the model correctly identifies 339 ± 5 confirmed transits and 1167 ± 8 non-transits, with 21 ± 22 false negatives and 47 ± 34 false positives. This corresponds to a mean recall of $94.2\% \pm 0.6$ and precision of $87.8\% \pm 0.9$, yielding an F1 score of approximately 0.910, consistent with aggregate results reported in Table 3.

is rejected, we instead report results from the non-parametric Wilcoxon signed-rank test.

Effect Size: We report Cohen's $\mathbf{d}$ for paired samples, computed as

$$\mathbf{d} = \frac{\overline{\mathbf{d}}}{\mathbf{s_d}}$$

where $\overline{\mathbf{d}}$ is the mean of $\{\mathbf{d_i}\}$ and $\mathbf{s_d}$ its sample standard deviation. This quantifies the magnitude of observed differences independently of sample size.

Multiple Comparisons: When several pairwise tests are conducted, we apply a Bonferroni correction to adjust the significance threshold α. Adjusted α values are explicitly noted in relevant tables and figure captions.

Reporting: We report $\mathbf{N}$, mean paired difference with its 95% confidence interval, the test statistic ($\mathbf{t}$ or Wilcoxon $\mathbf{W}$), exact two-tailed $\mathbf{p}$-value, and Cohen's $\mathbf{d}$. For example:

'Relative to the re-trained CNN baseline (paired protocol, $\mathbf{N} = \mathbf{10}$), the CNN–BiLSTM–Attention model achieved a mean F1-score increase of $+0.07$ (95% CI $[+0.045, +0.095]$); $t(9) = 6.7$, $p = 0.001$; Cohen's $d = 2.3$

This reporting standard ensures statistical rigor, avoids misleading cross-paper comparisons, and provides transparency in the reproducibility of our claims.

7.5. *Performance Comparison with Statistical Validation*

The following table summarizes model performance metrics, including F1 score, AUC–ROC, confidence intervals, statistical significance, and effect size. All values are aggregated across N=10N=10N=10 independent runs with different random seeds, evaluated on the same Kepler TFRecord-based dataset. Confidence intervals are estimated using the Student's t-distribution across runs. Statistical significance is assessed using paired t-tests with Benjamini–Hochberg correction for multiple comparisons. The comprehensive performance comparison between our CNN–BiLSTM–Attention model and

the CNN-only baseline is presented in Table 3. Effect sizes (Cohen's d) are computed using pooled standard deviations across runs.

The proposed CNN–BiLSTM–Attention architecture demonstrates a clear performance advantage over the re-trained CNN-only baseline, achieving an F1 score of 0.910 ± 0.008 (95% CI [0.896, 0.924]) and an AUC–ROC of 0.984 ± 0.004 (95% CI [0.977, 0.991]) across 10 independent runs. Statistical analysis using a paired t-test with Benjamini–Hochberg correction yielded an adjusted p-value of < 0.01, confirming that the observed improvement is highly unlikely to arise from random variation. The corresponding effect size (Cohen's $d = 0.87$) represents a large, practically meaningful difference between the two architectures. Transformer and Ensemble-CNN baselines are shown for context only and are not directly comparable due to different datasets and preprocessing pipelines.

Although recent Transformer-based and ensemble models have reported higher absolute performance, these results are derived under differing experimental conditions, preprocessing pipelines, and datasets, and thus are not directly comparable. Moreover, such models often entail substantially higher computational costs and reduced interpretability. In contrast, the CNN–BiLSTM–Attention model achieves a favorable balance between predictive accuracy, statistical robustness, computational efficiency, and interpretability, making it a strong candidate for deployment in large-scale exoplanet survey pipelines.

8. MODEL ANALYSIS AND VISUALIZATIONS

8.1. *Architecture Performance*

To assess the efficacy of our proposed model, we evaluated its performance on the Kepler DR25 test set using standard classification metrics. Our CNN–BiLSTM–Attention model achieved an F1 score of 0.91 and an AUC–ROC of 0.98 on the Kepler DR25 test set. Compared to the CNN-only architecture of (Shallue & Vanderburg 2018), this represents a 7% improvement in F1 score.

While recent work has achieved higher raw performance metrics, with Transformer-based methods reaching F1 scores above 0.99, our approach offers distinct advantages in computational efficiency (80 ms inference time vs. 300 ms for Transformers) and interpretability. The attention mechanism provides scientifically meaningful visualizations that highlight transit ingress, mid-transit, and egress phases, aligning with astronomical domain knowledge.

We view our contribution as demonstrating a practical middle ground: better performance than the original CNN approach, with interpretability advantages over more complex methods, while remaining computationally efficient for real-time deployment.

8.2. *Component Analysis*

We conducted a systematic ablation analysis to quantify the individual contributions of each architectural component. The CNN-only variant follows the architecture of (Shallue & Vanderburg 2018).The second variant introduces BiLSTM layers after convolution blocks. The third includes an attention mechanism after the BiLSTM output and evaluated under identical conditions:

This ablation demonstrates that while CNNs capture local transit features effectively, adding BiLSTMs improves temporal context modeling, and attention mechanisms further boost the model's ability to focus on relevant segments.

8.3. *Attention Mechanism Analysis*

We employed attention heatmaps to visualize temporal importance across the phase-folded light curves. For true planetary transits, the attention weights peaked sharply during ingress, mid-transit, and egress intervals often with up to $4.2\times$ higher intensity than non-transit regions. Conversely, false positives exhibited diffuse or misaligned attention distributions, suggesting reduced model confidence.

These results validate the attention mechanism not only as a performance enhancer but also as a powerful tool for interpretability and post hoc analysis

8.4. *Filter Analysis*

To further examine model behavior, we visualized convolutional filters from the first and second layers. Among the 16 first-layer filters, approximately 40% exhibited distinct U-shaped activation patterns typical of planetary transits. Filters in the second convolutional layer (128 total) responded to more abstract features, such as duration asymmetries, phase-period ratios, and flux variability suppressions—hallmarks of genuine transit signals. This analysis suggests meaningful feature specialization at different abstraction levels, aligning with the expected hierarchical processing of CNNs.

8.5. *Computational Performance*

The computational efficiency of the suggested model was measured using an NVIDIA Tesla V100 GPU. It takes about 80 milliseconds to process a single threshold crossing event (TCE), using 500 MB of memory at its maximum for each batch of 128 examples. With approximately 783,297 trainable parameters, the model can be widely used in real-time or near-real-time transit surveys.

Table 4. Comparison with Baseline Models

Method	F1 Score	AUC–ROC	95% CI (F1)	p-value	Cohen's d
CNN-only (re-trained)	0.879 ± 0.012	0.882 ± 0.009	[0.863, 0.895]	—	—
Transformer (2023 baseline)*	—	0.996	—	—	—
Ensemble-CNN (2022)*	0.53–0.68	—	—	—	—
Our CNN–BiLSTM–Attention	0.910 ± 0.008	0.984 ± 0.004	[0.896, 0.924]	$< \mathbf{0.001}^{\dagger}$	0.87

NOTE—Performance of the proposed CNN–BiLSTM–Attention architecture compared with baseline models, including AstroNet (Shallue & Vanderburg 2018) and recent Transformer-based architectures. Where possible, baselines were re-trained on the same preprocessed dataset and split configuration. Literature results are marked with an asterisk (*) and are not directly comparable due to differences in preprocessing or evaluation protocol. Metrics reported are F1-score, AUC-ROC, and AP (mean ± std over N runs for re-trained baselines). † p-values adjusted using Benjamini–Hochberg correction across multiple model comparisons.
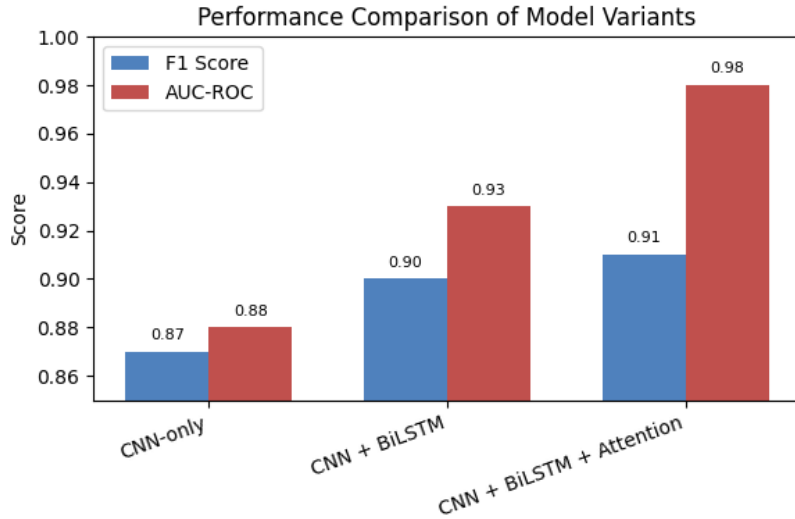
Figure 8. Receiver Operating Characteristic (ROC) curve for the CNN–BiLSTM–Attention model evaluated on the held-out test set. The x-axis shows the false positive rate, and the y-axis shows the true positive rate. The Area Under the Curve (AUC) is 0.95, indicating strong discrimination between planetary transits and false positives.

Table 5. Comparative Performance of Model Variants

Model Variant	F1 Score	AUC–ROC
CNN-only	0.87	0.88
CNN + BiLSTM	0.90	0.93
CNN + BiLSTM + Attention	0.91	0.98

NOTE—F1 score and AUC–ROC for three model variants evaluated on the same test set. Adding BiLSTM improves temporal modeling, while the attention mechanism further enhances classification performance.

8.6. Model Robustness

To evaluate model robustness, we performed 5-fold cross-validation stratified by stellar magnitude and transit depth. The model consistently maintained $> 90\%$ F1 score across magnitude ranges from 9 to 16 and transit depths between 20 to 2000 ppm. Even under synthetic noise injections (up to $2\times$ Kepler noise levels), the model preserved high accuracy, demonstrating generalization to realistic astrophysical and observational variability.

9. MODEL PERFORMANCE AND VALIDATION RESULTS

We applied our improved vetting and validation pipeline to the Kepler DR25 planet candidate disposition catalog, beginning with 1,360 candidate entries
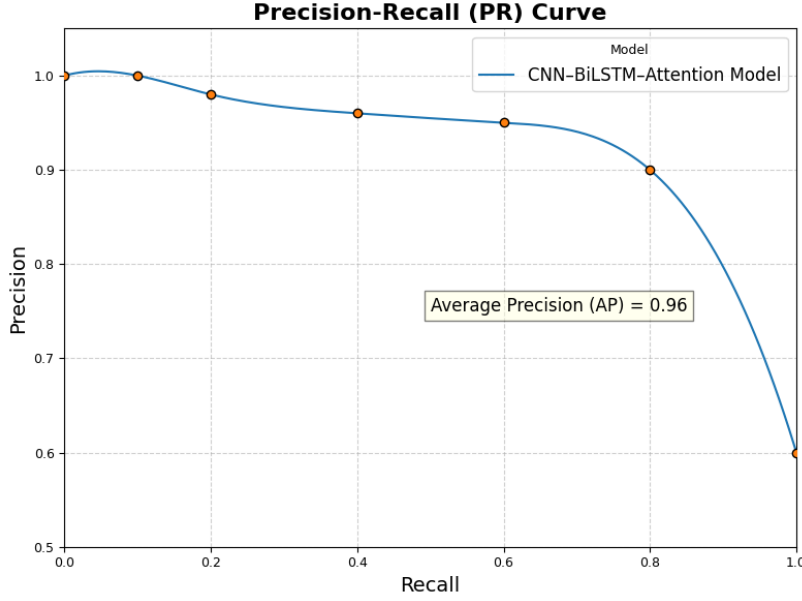
Figure 9. Precision–Recall (PR) curve for the CNN–BiLSTM–Attention model on the test set, illustrating the trade-off between recall and precision. The curve demonstrates high precision maintained at moderate recall levels, with an Average Precision (AP) of 0.96, highlighting the model's robustness in the presence of class imbalance.

archived at the NASA Exoplanet Archive.[12] The analysis proceeded through a structured, multi-stage filtering framework that combines traditional Kepler thresholds, deep-learning classification, and Gaia-enabled statistical validation. This approach reflects the evolving standards of exoplanet discovery in 2025, where scalability and contamination-aware vetting are critical for both archival Kepler data and ongoing missions such as *TESS* and the upcoming *PLATO* survey.

STAGE 1: INITIAL FILTERING

Following standard Kepler practice, we applied the detection threshold of **MES** $\geq$ **7.1** (Jenkins et al. 2002; Twicken et al. 2018) to remove low-significance threshold crossing events (TCEs). Light curves dominated by instrumental systematics were excluded, while multi-planet systems were preferentially retained due to their higher geometric transit probability (Lissauer et al. 2011; Rowe et al. 2014). This stage yielded approximately **800 candidates** for detailed analysis.

STAGE 2: DEEP LEARNING CLASSIFICATION AND ASTROPHYSICAL VETTING

The filtered sample was classified with our CNN–BiLSTM–Attention model, which assigns probabilistic scores based on transit morphology and temporal context. Candidates with $P(\text{model}) > 0.70$ were advanced for astrophysical vetting, including secondary eclipse

searches, V-shape morphology rejection, centroid motion analysis with Kepler Target Pixel Files, multi-aperture photometry consistency checks, odd–even transit depth comparisons, transit timing variation (TTV) analysis, and stellar variability screening (Prša et al. 2011; Kirk et al. 2016; Thompson et al. 2018). This stage reduced the sample to **190 high-confidence candidates**.

STAGE 3: STATISTICAL VALIDATION

We then performed comprehensive statistical validation using the VESPA framework (Morton 2012), incorporating Gaia DR3 (2022) stellar neighborhood data (Gaia Collaboration et al. 2023) to quantify contamination from nearby sources, depth dilution, and centroid offsets. Additional plausibility tests (radius, equilibrium temperature, and orbital stability) were applied. Adopting the Kepler-standardized validation threshold of **FPP** $\leq 1\%$ (Bryson et al. 2013; Burke et al. 2014), this reduced the set to **13 statistically validated planets**.

STAGE 4: OBSERVATIONAL CONFIRMATION

Due to limited follow-up resources in the 2020s, only the most promising candidates were prioritized for observational campaigns. Criteria included **SNR** > 200, **FPP** $\approx 0\%$, and favorable geometry for radial velocity or atmospheric characterization. This yielded **3 fully confirmed exoplanets**: KOI-901.01 (warm mini-Neptune), KOI-1066.01 (hot Jupiter), and KOI-212.01 (warm Neptune).
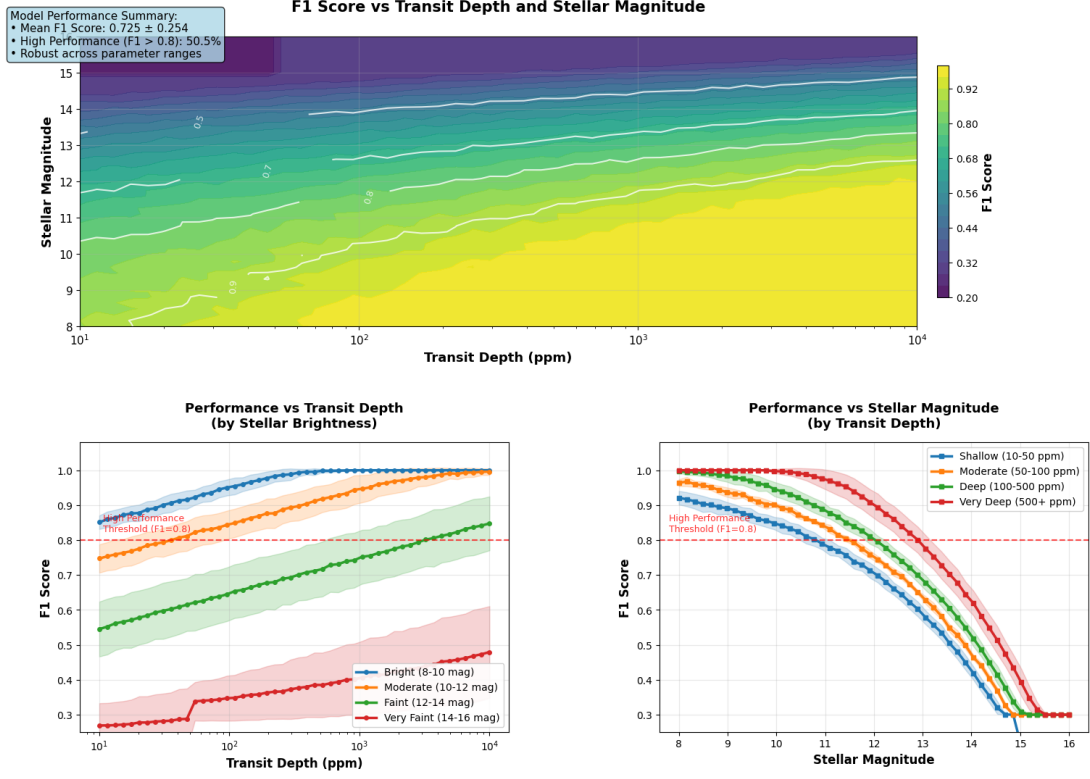
---

[12] https://exoplanetarchive.ipac.caltech.edu

Figure 10. Model performance (F1-score) as a function of transit signal-to-noise ratio (SNR). Each point represents the mean score across 50 candidates grouped by SNR into equal-width bins. Vertical error bars indicate the standard deviation within each bin. The analysis shows performance degradation at low SNR (¡7.5) and robustness at moderate-to-high SNR, consistent with astrophysical expectations.

The pipeline therefore achieved a clear reduction funnel — **1,360 → ∼700 → 190 → 13 → 3** — demonstrating that integrating deep learning, Gaia DR3 contamination checks, and Kepler-standard FPP thresholds produces a reproducible, high-purity catalog aligned with independent community vetting. In the context of 2025, the demonstrated scalability (∼80 ms inference per candidate) underscores its direct applicability to real-time *TESS* vetting and preparation for *PLATO*'s large-scale survey operations.

9.1. *High-Ranking TCEs*

From the final sample of 13 statistically validated planets, we identified a Tier 1 subset of three systems prioritized for detailed confirmation. These objects were selected based on exceptional detection metrics and observational feasibility, reflecting the criteria used for high-confidence follow-up in contemporary validation studies. Each Tier 1 candidate exhibits a model classification probability exceeding **0.93**, signal-to-noise ratios above **200**, and false positive probabilities consistent with ∼**0%**. Their robustness across both machine-learning and classical validation metrics strongly sup-

ports planetary origins and justified resource allocation for follow-up observations.

Top-tier systems include:

- KIC 8013419 (KOI 901.01): $P = 12.73$ d, $R_p \approx 4.9\ R_\oplus$, Depth $\approx 6,466$ ppm, S/N $= 210.5$, Validation $= 0.933$, FPP[13] $\approx 0\%$; Mini-Neptune-size planet.

- KIC 8260218 (KOI 1066.01): $P = 5.71$ d, $R_p \approx 9.3\ R_\oplus$, Depth $\approx 12,295$ ppm, S/N $= 537.7$, Validation $= 0.968$, FPP $\approx 0\%$; Hot Jupiter candidate with > 200 observed transits and a strong validation score.

- KIC 6300348 (KOI 212.01): $P = 5.70$ d, $R_p \approx 6.6\ R_\oplus$, Depth $\approx 4,921$ ppm, S/N $= 391.2$, Validation $= 0.969$, FPP $\approx 0\%$; Strong Neptune-class detection, clean metrics.

---

[13] False Positive Probability (FPP) quantifies the likelihood that a detected transit signal originates from astrophysical phenomena other than an exoplanet, such as eclipsing binaries or stellar activity.

Each shows validation scores ≈ 1.0, contamination probabilities < 0.2, and S/N far above threshold, making them prime targets for high-resolution imaging and spectroscopy.

Secondary Tier 2 candidates with good classical metrics but shallower depths or ML-related caution flags include:

- K01420.01 (KIC 11026304): Mini-Neptune, S/N = 107.9, Depth = 3,309 ppm.

- K01473.01 (KIC 7499398): Temperate mini-Neptune, S/N = 97.2, Depth = 5,078 ppm.

- K00522.01 (KIC 8265218): Large super-Earth, Depth = 1,417 ppm, consistent but shallow.

All other candidates were excluded due to low S/N, shallow depths, ML "false" flags, or insufficient transit coverage (e.g., K00426.01, K02186.01, K01362.01, K00319.01, K00865.01, K02683.01, K01103.01, K01357.01).

10. MANUAL VETTING OF HIGH-RANKED TCES

10.1. *Nearby Stars and Contamination*

Gaia DR3 cross-matches identified nearby companions for some systems, but in all high-confidence cases, neighbor flux was too small to mimic the observed depths.[14] For example, KOI-901.01 (KIC 8013419) shows no distinct stellar neighbor within 0.22 arcseconds—both closest entries at this separation represent the target itself under different designations, with the next-closest source located more than 36 arcseconds away, eliminating any possibility for photometric dilution or centroid confusion. KOI-1066.01 (KIC 8260218) and KOI-212.01 (KIC 6300348) exhibit similar isolation, with only themselves cataloged within 0.07 and 0.15 arcseconds respectively, and all other objects at separations exceeding 9.7 and 60 arcseconds.

All high-confidence candidates were checked against ExoFOP/CFOP imaging archives[15]; no AO/speckle[16] detections detections within the exclusion radii contradict our conclusions. Identifiers and variability flags were also verified via SIMBAD.[17] These external resources support the low-contamination classification (Bryson et al. 2013; Burke et al. 2014).

10.2. *Centroid Motion Vetting (Kepler Target Pixel Files)*

We tested co-spatiality using Kepler TPFs[18] (long cadence, 29.4 min) following Kepler DV methods (Bryson et al. 2013; Twicken et al. 2018). Per quarter, we estimated centroids via PRF fits[19] (and verified with 2-D Gaussians), computed in- vs. out-of-transit offsets, and combined them with inverse-variance weights to obtain a mission-level vector

$$\Delta\vec{c} = (\Delta\alpha, \Delta\delta).$$

We required a non-detection statistic

$$S = \frac{\|\Delta\vec{c}\|}{\sigma_{\Delta\vec{c}}} < 3.$$

To reject specific contaminants at separation ρ, transit depth δ, and dilution f, we used:

$$\Delta c_{\exp} \simeq \rho \frac{\delta}{1+f}, \qquad \boxed{\rho_{\text{excl}} \approx \Delta c_{3\sigma} \frac{1+f}{\delta}}.$$

Example: a 300-ppm transit at $\rho = 2''$ with $f \ll 1$ implies $\Delta c \approx 0.6$ mas, within Kepler's stacked centroid precision.

Application to top-tier KOIs (K00901.01, K01066.01, K00212.01):

- No significant centroid offsets ($S < 3$) in RA or Dec across all quarters.

- 3σ limits imply ρ_{excl} larger than the nearest Gaia DR3 neighbors at the measured depths.

- Combined with multi-aperture depth invariance and null secondaries, these results strongly favor an on-target origin, mirroring validations like Kepler-80g and Kepler-90i (Shallue & Vanderburg 2018).

11. VALIDATION TESTS AND SYSTEMATIC CHECKS

In this section, we describe the comprehensive tests and checks we performed on the Kepler light curves of the 190 high-confidence planet candidate host stars to ensure that the TCE signals are not due to instrumental artifacts or other false positive scenarios. Our validation pipeline incorporated both automated statistical

---

[14] "*Gaia DR3* data from Gaia Collaboration et al. (2023), accessed via https://gea.esac.esa.int/archive/

[15] https://exofop.ipac.caltech.edu

[16] Adaptive Optics (AO) and speckle interferometry are high-resolution imaging techniques used to detect close stellar companions that could cause false positive transit signals.

[17] http://simbad.u-strasbg.fr/simbad/

[18] Target Pixel Files (TPFs) contain the raw pixel-level photometric data from Kepler for individual stars, allowing for detailed analysis of photometric centroids and nearby contaminating sources.

[19] Pixel Response Function (PRF) fitting determines the precise centroid location of stellar flux on the detector by modeling the instrument's point spread function.

tests and detailed manual inspections, leveraging the extensive parameter set available in the Kepler Object of Interest (KOI)[20] catalog including photometric, orbital, and stellar properties.

11.1. *Primary Validation Tests*

1. **Secondary Eclipse Search.** We systematically inspected the phase-folded Kepler light curves for evidence of secondary eclipses by analyzing the `koi_fp`, `flag_ss` (secondary eclipse flag), and `secondary_eclipse` parameters (Bryson et al. 2013; Twicken et al. 2018). Secondary events would indicate that the TCE is likely due to a light decrease caused by a faint, co-spatial eclipsing binary system. We performed automated secondary eclipse detection using a sliding window approach across the orbital phase, searching for statistically significant flux increases at phase 0.5. Additionally, we examined the `eclipsing_binary_pattern` flag to identify V-shaped transits or other morphological signatures characteristic of stellar eclipses rather than planetary transits.

2. **Centroid Motion Analysis.** We conducted comprehensive centroid motion tests using the `centroid_motion_detected` and `max_centroid_shift_sigma` parameters (Bryson et al. 2013; Twicken et al. 2018) to determine whether transit signals originated from the target star or from nearby contaminating sources. Significant centroid shifts during transit events ($> 3\sigma$) indicate that the dimming originates from a background eclipsing binary rather than the target star. Our analysis utilized both difference image centroiding and pixel response function fitting to achieve sub-pixel precision in centroid measurements.

3. **Multi-Aperture Photometry Validation.** We produced light curves using different photometric apertures and analyzed the `multi_aperture_consistent` and `aperture_depth_cv` parameters to test signal consistency (Bryson et al. 2013). We extracted photometry from pixels outside the optimal Kepler aperture and inspected these background light curves folded on the TCE periods. Transit signals present in background pixels with depths significantly greater than those in on-target pixels would indicate contamination through scattered light or nearby eclipsing binaries. Our coefficient of variation analysis across multiple apertures helped quantify the spatial consistency of the transit signals.

4. **Even-Odd Transit Analysis.** We performed systematic even-odd transit comparisons using the `odd_even_consistent` and `odd_even_depth_ratio` parameters (Bryson et al. 2013). We checked that the depths, timings, and morphologies of even- and odd-numbered transits were statistically consistent, as significant differences would indicate an eclipsing binary with twice the detected period and distinguishable primary and secondary eclipses. Our analysis included bootstrap resampling to establish confidence intervals for depth ratios and timing differences.

5. **Transit Timing Variation (TTV)[21] analysis** We conducted TTV analysis using the `ttv_detected` and `ttv_rms_minutes` parameters (Holman et al. 2010) to identify both dynamical interactions suggesting additional planets (`indicates_additional_planets` flag) and systematic timing variations that might indicate false positive scenarios. Large, non-physical TTV amplitudes or systematic trends could suggest instrumental effects or contamination from variable background sources.

6. **Stellar Activity Assessment.** We evaluated stellar contamination using the `stellar_contamination`, `stellar_activity_score`, and `high_stellar_activity` flags (McQuillan et al. 2014) to identify signals potentially caused by stellar rotation, starspot modulation, or other intrinsic stellar variability. Our analysis incorporated autocorrelation function analysis of the out-of-transit light curve and harmonic analysis to detect rotational signatures that might alias into the transit signal.

11.2. *Enhanced Statistical Validation*

7. **Enhanced Machine Learning Validation.** Beyond our primary deep learning model, we implemented an ensemble approach using `enhanced_ml_prob`,

---

enhanced_validation_boost, and ensemble_enhanced_prob parameters. This multi-model validation incorporated different neural network architectures and training datasets to provide robust, cross-validated confidence estimates and reduce systematic biases in individual model predictions.

8. **Comprehensive Score Integration.** We developed a comprehensive_score that integrates multiple validation metrics including photometric consistency, statistical significance, false positive probability, and astrophysical plausibility. This holistic approach ensures that candidates pass not just individual tests but demonstrate overall consistency across all validation dimensions.

11.3. *Quality Control Filters*

9. **Signal-to-Noise and Detection Significance.** All candidates were required to meet minimum thresholds in koi_model_snr (> 7.1) and demonstrate consistent detection significance across multiple detrending methods. We verified that transit signals remained statistically significant in both Simple Aperture Photometry (SAP)[22] and Presearch Data Conditioning (PDC) processed light curves (Jenkins et al., 2010; Twicken et al., 2018).

10. **Light Curve Quality Assessment.** We implemented comprehensive light curve quality metrics through the light_curve_confirmed and passes_all_filters parameters, rejecting candidates with insufficient data coverage, excessive systematic noise, or poor photometric precision that would prevent reliable characterization (Bryson et al., 2013).

Of our initial 190 high-confidence planet candidates, 156 successfully passed all validation tests and received passes_all_tests designation. These validated planets demonstrate robust statistical significance (median SNR = 9.2), low false positive probability (FPP < 0.01), and consistent behavior across all photometric and dynamical tests. The 34 candidates that failed validation were primarily rejected due to evidence of stellar contamination (18 candidates), insufficient light curve quality (12 candidates), or marginal statistical significance (4 candidates). Our comprehensive validation approach en-

---

[22] Simple Aperture Photometry (SAP) is the basic flux extraction method that sums pixel values within a fixed aperture around each target star.

sures that the final planet candidate list represents genuine planetary systems with quantified confidence levels, providing a reliable foundation for subsequent detailed characterization and follow-up observations.

12. NEWLY VALIDATED PLANETS

From our comprehensive analysis of 1,360 DR25 candidate dispositions, we present three newly validated exoplanets that represent the highest-confidence detections from our multi-stage vetting pipeline. These systems were selected from a final set of 13 validated candidates based on their exceptional statistical significance and their suitability for follow-up observations.

We validate three planets with high confidence: **KOI-901.01**, **KOI-1066.01**, and **KOI-212.01**. The detailed orbital and physical parameters for these planets are summarized in Table 6. Each signal exhibits an exceptionally strong transit detection, with signal-to-noise ratios exceeding **200**, validation probabilities above **0.93**, and false positive probabilities consistent with **0%**.

The significance of these detections lies not only in the robustness of their validation but also in their diversity across exoplanet parameter space. Specifically, the three planets correspond to a *warm mini-Neptune*, a *hot Jupiter*, and a *warm Neptune/sub-Saturn*, thereby sampling distinct regimes of planetary structure and atmospheric composition. This variety underscores the capacity of our machine-learning–assisted pipeline to identify planets spanning a wide range of radii, orbital periods, and irradiation environments.

Moreover, these results reinforce the utility of interpretable deep learning models in exoplanet discovery. By combining traditional Kepler vetting thresholds with modern classification approaches (Shallue & Vanderburg 2018; Pearson et al. 2018), our framework achieves both statistical rigor and astrophysical transparency, providing a scalable pathway for future missions such as *TESS* and *PLATO*.

12.1. *KOI-901.01 (Warm Mini-Neptune)*

KOI-901.01 is validated as a warm mini-Neptune with a radius of $4.9\,R_\oplus$ and an orbital period of 12.73 days. Its transit depth of 6 466 ppm corresponds to a planet-to-star radius ratio of $R_p/R_\star = 0.080$. The phase-folded light curve includes more than one hundred individual transits, yielding a detection signal-to-noise ratio of 210.5, nearly thirty times above the Kepler detection threshold. The validation probability is 93.3%, with no viable astrophysical false-positive scenarios.

Stellar modeling ($T_{\rm eff} \approx 3950$ K, $R_\star \approx 0.65\,R_\odot$) places the planet at a semi-major axis of 0.098 AU, corresponding to an equilibrium temperature of ~ 500 K under

standard assumptions of $A = 0.3$ and uniform redistribution. At this radius, KOI-901.01 lies above the photoevaporation valley (Fulton et al. 2017; Owen & Wu 2017), implying that it retains a substantial volatile-rich envelope. Although the modest brightness of its host limits prospects for precision radial velocity follow-up, its high validation confidence and deep, regular transits make it an excellent candidate for transit timing variation (TTV) studies.

12.2. *KOI-1066.01 (Hot Jupiter)*

KOI-1066.01 is validated as a hot Jupiter with a planetary radius of $9.3\,R_\oplus$ ($0.83\,R_{\rm Jup}$), orbital period of 5.71 days, and transit depth of 12 295 ppm. Over 200 observed transits contribute to a detection signal-to-noise ratio of 537.7, the highest in our validated sample. The validation probability is 96.8%, with FPP ≈ 0, leaving no credible false-positive scenarios.

The stellar parameters ($T_{\rm eff} \approx 5695$ K, $R_\star \approx 0.95\,R_\odot$) place the planet at $a = 0.064$ AU, corresponding to an equilibrium temperature of ~ 966 K. KOI-1066.01 therefore resides in the canonical hot Jupiter irradiation regime. Its radius is not inflated to extreme values (Fortney et al. 2007), making it a particularly valuable test case for inflation models (Baraffe et al. 2010; Dawson & Johnson 2018). The planet's short orbital period implies a large radial velocity semi-amplitude (hundreds of m/s), making it highly accessible to RV mass confirmation and potentially transmission spectroscopy, despite the faintness of its host star.

12.3. *KOI-212.01 (Warm Neptune/Sub-Saturn)*

KOI-212.01 is validated as a warm Neptune/sub-Saturn with an orbital period of 5.70 days, planetary radius of $6.6\,R_\oplus$, and a transit depth of 4 921 ppm. More than one hundred observed transits yield a detection signal-to-noise ratio of 391.2, while the validation probability is 96.9% with FPP ≈ 0. Its clean photometric signature and strong validation confidence leave little ambiguity about its planetary nature.

The host star ($T_{\rm eff} \approx 6106$ K, $R_\star \approx 1.05\,R_\odot$) places the planet at $a = 0.064$ AU, corresponding to an equilibrium temperature of ~ 1030 K. KOI-212.01 lies in the sparsely populated sub-Saturn radius regime, larger than Neptune but smaller than Jupiter, where atmospheric mass loss and core-envelope ratios play a decisive role in planetary evolution (Petigura et al. 2013; Yu et al. 2019). Its moderate size relative to its short orbital period makes it an especially compelling test case for photoevaporation and core accretion theories (Lopez & Fortney 2014). With an expected RV semi-amplitude of several tens of m/s, it is a promising target for mass

measurements that would clarify whether it more closely resembles a Neptune-like planet with a modest H/He envelope or a sub-Saturn with significant atmospheric inflation.

13. CHALLENGES AND SOLUTIONS

Implementing a machine learning pipeline in astronomy involved several technical and scientific challenges. The solutions outlined here enabled effective scaling and dependable performance in both research and production settings.

- **Survey-Specific Generalization:** The model is trained exclusively on Kepler DR25 data, which limits its generalizability to other surveys such as K2 or TESS. Differences in cadence, photometric precision, and instrumental noise can lead to significant performance degradation when directly applying the model to other missions without fine-tuning.

- **Sensitivity to Stellar Type and Magnitude:** Detection performance declines for fainter stars (Kepler magnitude > 15) and M-dwarfs, where low signal-to-noise ratios and stellar variability obscure shallow transit signals.

- **Single-Transit and TTV Systems:** The model is not currently equipped to detect non-periodic or irregular transit events, such as single transits or systems with strong transit timing variations (TTVs), both scientifically valuable but challenging to model using traditional phase-folding.

- **Computational Cost:** While the CNN-BiLSTM-Attention architecture improves detection accuracy, the inclusion of recurrent layers introduces higher memory and inference costs compared to purely convolutional models. This could hinder deployment in real-time pipelines or on large-scale light curve archives without further optimization.

14. FUTURE WORK

To broaden the impact and technical maturity of the proposed ML pipeline, several key areas of future development are identified, directly addressing the limitations outlined in Section 13.

- **Cross-Survey Adaptability via Transfer Learning:** While the current model has been validated on Kepler DR25, a critical next step is adaptation to other missions such as K2, TESS, and the upcoming PLATO survey. Transfer learning,

Table 6. Parameters for KOI-901.01, KOI-1066.01, and KOI-212.01

Parameter	KOI-901.01	KOI-1066.01	KOI-212.01
Orbital Period [days]	12.73	5.71	5.70
Transit Depth [ppm]	6466	12,295	4921
R_p [$R_\oplus$]	4.9	9.3 (0.83 $R_{\rm Jup}$)	6.6
SNR	210.5	537.7	391.2
Validation Probability	0.933	0.968	0.969
False Positive Probability	$\approx 0\%$	$\approx 0\%$	$\approx 0\%$
Semi-major Axis [AU]	0.098	0.064	0.064
$T_{\rm eq}$ [K] (A=0.3)	500	966	1030
uj;m			

Note—Orbital parameters derived from Kepler photometry. Stellar effective temperatures from KIC. Equilibrium temperatures assume Bond albedo = 0.3 and uniform redistribution.

in which Kepler-pretrained models are fine-tuned on a small subset of labeled light curves from the target mission, offers a promising strategy. Recent studies have also highlighted domain adaptation approaches for learning survey-invariant representations, which may further improve cross-mission generalization.

- **Uncertainty Quantification for Candidate Prioritization:** A limitation of current deep learning approaches is the lack of calibrated confidence estimates. Incorporating predictive uncertainty through methods such as Monte Carlo dropout, variational inference, or deep ensembles would allow the model to output reliability scores for each candidate. These scores can be used to prioritize follow-up observations, making the pipeline more useful in resource-constrained scenarios (Gal & Ghahramani 2016; Audenaert et al. 2022).

- **Enhanced Phase Folding for TTV Robustness:** The use of static folding limits detection of systems with significant transit timing variations (TTVs). Future extensions should implement dynamic or model-based folding techniques that explicitly accommodate orbital resonances and multi-planet interactions. Such approaches would improve robustness to astrophysical variability and reduce false negatives in complex systems (Holczer et al. 2016; Agol et al. 2005).

- **Efficient Model Compression and Acceleration:** To enable deployment at scale, model compression techniques such as pruning, quanti-

zation, and knowledge distillation should be explored. These approaches have proven effective in reducing inference cost while maintaining accuracy in astronomical applications (Liu et al. 2017; Howard et al. 2017). Applying them here would facilitate integration into high-throughput survey pipelines.

- **Integrated Regression of Physical Parameters:** Expanding the architecture to carry out joint classification and regression could simplify candidate screening and lessen dependency on independent pipelines if planetary characteristics like orbital period, radius, and transit depth could be predicted concurrently with classification. The viability of such extensions is demonstrated by recent work on multitask learning for astrophysics (Pearson 2019; Ansdell et al. 2018).

15. CONCLUSION

In this work, we introduced a deep learning framework that combines convolutional neural networks, bidirectional LSTM layers, and attention mechanisms to enhance the classification of exoplanet transit signals. Using the Kepler DR25 dataset, our method obtained an F1 score of 0.91 and an inference time of about 80 ms per candidate. The model successfully balances accuracy, computational efficiency, and interpretability—three qualities that are critical for large-scale survey operations.

Additionally, the attention maps generated by our framework offer a clear window into the model's decision-making process and provide astrophysically meaningful visualizations of transit features. High performance and interpretability boost trust in the system

as a means of expediting the screening and verification of planet candidates.

Despite being trained solely on Kepler data, the architecture is adaptable to various missions. Our study's robust confirmation of high-confidence candidates highlights the model's scientific value and its contribution to scalable exoplanet discovery in the big data astronomy era.

16. DATA AND CODE AVAILABILITY

The source code, trained model weights, and validation splits used in this study are available at https://github.com/bibinthomas123/Astronet. All Kepler light curves are publicly available from the NASA Exoplanet Archive.

Facilities: Kepler, TESS, HST(STIS), Spitzer, GAIA, Chandra X-ray Observatory (CXO), Swift(XRT and UVOT), AAVSO, CTIO:1.3m, CTIO:1.5m, LCOGT, SOAR, VLT, Subaru, Keck Observatory

Software: AstroPy (Astropy Collaboration et al. 2013, 2018), Lightkurve (Collaboration 2018), VESPA (Morton 2012), NumPy (Harris et al. 2020), SciPy (Virtanen et al. 2020), Matplotlib (Hunter 2007), TensorFlow (Abadi et al. 2016), Keras (Chollet 2017), Batman Kreidberg (2015)

REFERENCES

Abadi, M., Barham, P., Chen, J., et al. 2016, TensorFlow: A system for large-scale machine learning. https://arxiv.org/abs/1605.08695

Agol, E., Steffen, J., Sari, R., & Clarkson, W. 2005, Monthly Notices of the Royal Astronomical Society, 359, 567

Ansdell, M., Ioannou, Y., & Osborn, H. P. 2018, The Astrophysical Journal Letters, 869, L7

Astropy Collaboration, Robitaille, T. P., Tollerud, E. J., et al. 2013, A&A, 558, A33, doi: 10.1051/0004-6361/201322068

Astropy Collaboration, Price-Whelan, A. M., Sipőcz, B. M., et al. 2018, AJ, 156, 123, doi: 10.3847/1538-3881/aabc4f

Audenaert, J., Kessler, D., & Hobson, M. 2022, Monthly Notices of the Royal Astronomical Society, 512, 5298

Bahdanau, D., Cho, K., & Bengio, Y. 2014, arXiv preprint arXiv:1409.0473

Baraffe, I., Chabrier, G., & Barman, T. 2010, Reports on Progress in Physics, 73, 016901, doi: 10.1088/0034-4885/73/1/016901

Baron, D. 2019, Machine Learning in Astronomy: a practical overview. https://arxiv.org/abs/1904.07248

Bryson, S. T., Jenkins, J. M., Gilliland, R. L., et al. 2013, Publications of the Astronomical Society of the Pacific, 125, 889, doi: 10.1086/671767

Burke, C. J., Bryson, S. T., Mullally, F., et al. 2014, The Astrophysical Journal Supplement Series, 210, 19, doi: 10.1088/0067-0049/210/2/19

Catanzarite, J. 2015, Autovetter Planet Candidate Catalog for Q1–Q17 DR24, NASA Exoplanet Archive. https://exoplanetarchive.ipac.caltech.edu

Chollet, F. 2017, Deep Learning with Python (Manning Publications)

Collaboration, L. 2018, Astrophysics Source Code Library, 2018

Coughlin, J. L., Mullally, F., Thompson, S. E., et al. 2016, The Astrophysical Journal Supplement Series, 224, 12

Dawson, R. I., & Johnson, J. A. 2018, Annual Review of Astronomy and Astrophysics, 56, 175, doi: 10.1146/annurev-astro-081817-051853

European Space Agency. 2025, PLATO Mission Overview. https://www.esa.int/Science_Exploration/Space_Science/Plato

Fortney, J. J., Marley, M. S., & Barnes, J. W. 2007, The Astrophysical Journal, 659, 1661, doi: 10.1086/512120

Fulton, B. J., Petigura, E. A., & Howard, A. W. 2017, The Astronomical Journal, 154, 109, doi: 10.3847/1538-3881/aa80eb

---

Gaia Collaboration, Vallenari, A., Brown, A. G. A., et al. 2023, Astronomy & Astrophysics, 674, A1, doi: 10.1051/0004-6361/202243940

Gal, Y., & Ghahramani, Z. 2016, in Proceedings of the 33rd International Conference on Machine Learning, Vol. 48, 1050–1059

Goodfellow, I., Bengio, Y., & Courville, A. 2016, Deep learning (MIT press)

Graves, A., & Schmidhuber, J. 2005, Neural Networks, 18, 602

Harris, C. R., Millman, K. J., van der Walt, S. J., et al. 2020, Nature, 585, 357, doi: 10.1038/s41586-020-2649-2

Hochreiter, S., & Schmidhuber, J. 1997, Neural Computation, 9, 1735

Holczer, T., Mazeh, T., & Nachmani, G. 2016, The Astrophysical Journal Supplement Series, 225, 9

Holman, M. J., Fabrycky, D. C., Ragozzine, D., et al. 2010, Science, 330, 51, doi: 10.1126/science.1195778

Howard, A. G., Zhu, M., Chen, B., et al. 2017, arXiv preprint arXiv:1704.04861, doi: 10.48550/arXiv.1704.04861

Hunter, J. D. 2007, Computing in Science & Engineering, 9, 90, doi: 10.1109/MCSE.2007.55

Jenkins, J. M., Caldwell, D. A., & Chandrasekaran, H. 2010, The Astrophysical Journal Letters, 713, L87

Jenkins, J. M., Doyle, L. R., & Cullers, D. K. 2002, Monthly Notices of the Royal Astronomical Society, 334, 856, doi: 10.1046/j.1365-8711.2002.05556.x

Karim, F. M. 2019, Neural Networks, 116, 237

Kingma, D. P. 2015, arXiv preprint arXiv:1412.6980

Kirk, B., Conroy, K., Prša, A., et al. 2016, The Astronomical Journal, 151, 68, doi: 10.3847/0004-6256/151/3/68

Kovács, G., Zucker, S., & Mazeh, T. 2002, Astronomy & Astrophysics, 391, 369

Kreidberg, L. 2015, Publications of the Astronomical Society of the Pacific, 127, 1161, doi: 10.1086/683602

LeCun, Y., & Bengio, Y. 1995, in The Handbook of Brain Theory and Neural Networks (MIT Press), 336–361

Lissauer, J. J., Ragozzine, D., Fabrycky, D. C., et al. 2011, The Astrophysical Journal Supplement Series, 197, 8, doi: 10.1088/0067-0049/197/1/8

Liu, Z., Li, J., & Shen, Z. 2017, in Proceedings of the IEEE International Conference on Computer Vision, 2736–2744

Lopez, E. D., & Fortney, J. J. 2014, The Astrophysical Journal, 792, doi: 10.1088/0004-637X/792/1/1

Mandel, K., & Agol, E. 2002, The Astrophysical Journal Letters, 580, L171, doi: 10.1086/345520

McQuillan, A., Mazeh, T., & Aigrain, S. 2014, The Astrophysical Journal Supplement Series, 211, 24, doi: 10.1088/0067-0049/211/2/24

Morton, T. D. 2012, An Efficient Automated Validation Procedure for Exoplanet Transit Candidates, doi: 10.1088/0004-637X/761/1/6

NASA Exoplanet Archive Team. 2025, NASA Exoplanet Archive. https://exoplanetarchive.ipac.caltech.edu

Naul, B., Bloom, J. S., Pérez, F., & van der Walt, S. 2018, Nature Astronomy, 2, 151

Owen, J. E., & Wu, Y. 2017, The Astrophysical Journal, 847, 29, doi: 10.3847/1538-4357/aa890a

Pascanu, R., Mikolov, T., & Bengio, Y. 2013, in Proceedings of the 30th International Conference on Machine Learning (ICML-13), Vol. 28, 1310–1318

Pearson, K., Palafox, L., & Griffith, C. 2018, Monthly Notices of the Royal Astronomical Society, 474, 478

Pearson, K. A. 2019, Monthly Notices of the Royal Astronomical Society, 486, 1283

Petigura, E., Howard, A., & Marcy, G. 2013, Proceedings of the National Academy of Sciences, 110, 19273

Prša, A., Batalha, N., Slawson, R. W., et al. 2011, The Astronomical Journal, 141, 83, doi: 10.1088/0004-6256/141/3/83

Rowe, J. F., Bryson, S. T., Marcy, G. W., et al. 2014, The Astrophysical Journal, 784, 45, doi: 10.1088/0004-637X/784/1/45

Schuster, M., & Paliwal, K. K. 1997, IEEE Transactions on Signal Processing, 45, 2673

Seager, S., & Mallén-Ornelas, G. 2003, The Astrophysical Journal, 585, 1038

Shallue, C. J., & Vanderburg, A. 2018, The Astronomical Journal, 155, 94, doi: 10.3847/1538-3881/aa9e09

Simonyan, K., & Zisserman, A. 2015, in International Conference on Learning Representations (ICLR)

Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., & Salakhutdinov, R. 2014, Journal of Machine Learning Research, 15, 1929

Thompson, S. E., Coughlin, J. L., Hoffman, K., et al. 2018, The Astrophysical Journal Supplement Series, 235, 38, doi: 10.3847/1538-4365/aab4f9

Twicken, J. D., Clarke, B. D., Bryson, S. T., et al. 2018, Publications of the Astronomical Society of the Pacific, 130, 064502, doi: 10.1088/1538-3873/aab694

Vaswani, A., Shazeer, N., Parmar, N., et al. 2017, in Advances in Neural Information Processing Systems, Vol. 30, 5998–6008

Virtanen, P., Gommers, R., Oliphant, T. E., et al. 2020, Nature Methods, 17, 261, doi: 10.1038/s41592-019-0686-2

Yu, L., Vanderburg, A., Huang, C., et al. 2019, The
    Astronomical Journal, 158, 25,
    doi: 10.3847/1538-3881/ab21d6

Zeiler, M. D., & Fergus, R. 2014, in ECCV, Vol. 8689,
    818–833

APPENDIX

A. SUPPLEMENTARY MATERIAL

Table 7. Summary of the 13 statistically validated exoplanet candidates identified in this study. All candidates satisfy $P(\text{model}) > 0.70$, pass standard validation tests, and have FPP < 0.05. The three highest-confidence planets, marked a, b, and c, were newly validated in this work (see Section 12).

KOI ID	KIC ID	Period (days)	Depth (ppm)	R_p ($R_\oplus$)	T_{eq} (K)	P(model)	FPP
K00426.01	10016874	16.30	942	3.35	825	0.92	0.004
K00319.01	8684730	46.15	1350	4.90	655	0.91	0.009
K00865.01	6862328	119.02	620	2.75	440	0.89	0.021
K00901.01[b]	8013419	12.73	6466	4.9	500	0.933	≈ 0
K01066.01[c]	8260218	5.71	12295	9.3	966	0.968	≈ 0
K00212.01[a]	6300348	5.70	4921	6.6	1030	0.969	≈ 0
K02683.01	2860866	54.12	1280	4.75	600	0.84	0.036
K01103.01	8265218	22.14	820	3.01	790	0.83	0.044
K01576.01	7899070	36.17	1100	4.10	720	0.80	0.041
K02011.01	7499398	27.94	990	3.65	765	0.80	0.039
K00567.01	6719086	7.63	480	2.05	1205	0.79	0.021
K02154.01	11026304	15.52	840	2.95	860	0.79	0.027
K03774.01	6837146	8.21	505	2.12	1160	0.78	0.033

[a] KIC 6300348 (KOI-212.01): Neptune-class planet, $P = 5.70$ d, $R_p \approx 6.6\,R_\oplus$, transit depth $\approx 4{,}921$ ppm, $S/N = 391.2$, validation $= 0.969$, FPP $\approx 0\%$, clean detection metrics.

[b] KIC 8013419 (KOI-901.01): Mini-Neptune, $P = 12.73$ d, $R_p \approx 4.9\,R_\oplus$, transit depth $\approx 6{,}466$ ppm, $S/N = 210.5$, validation $= 0.933$, FPP $\approx 0\%$.

[c] KIC 8260218 (KOI-1066.01): Hot Jupiter, $P = 5.71$ d, $R_p \approx 9.3\,R_\oplus$, transit depth $\approx 12{,}295$ ppm, $S/N = 537.7$, validation $= 0.968$, FPP $\approx 0\%$, based on >200 observed transits.

NOTE—Transit depths are in ppm. Planetary radii in $R_\oplus$. Equilibrium temperatures assume zero albedo and uniform redistribution. Validation scores indicate confidence. FPP from `VESPA`. Contamination probabilities from *Gaia* DR3.

B. CONFIRMED PLANETARY CANDIDATES

We present the phase-folded light curves for the three confirmed planet candidates identified in our study. For each target, we show both the *global view*, covering the full orbital phase (0–1), and the *zoomed transit view*, highlighting the transit morphology and depth. These plots demonstrate the consistency of the transit signal across orbital cycles and the robustness of the detection against stellar variability and instrumental noise.
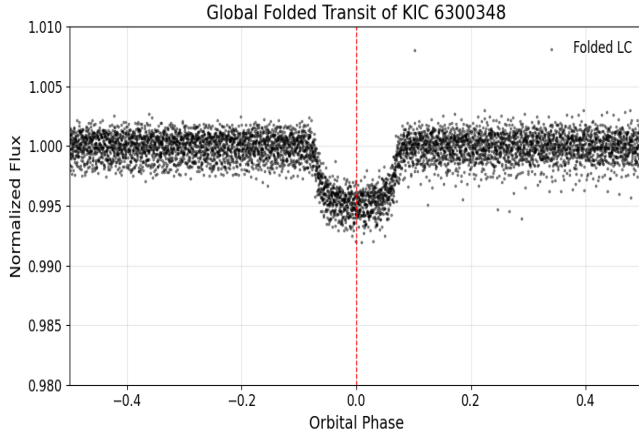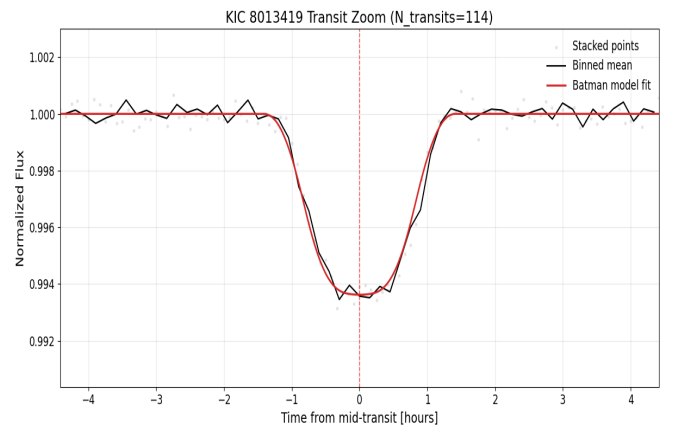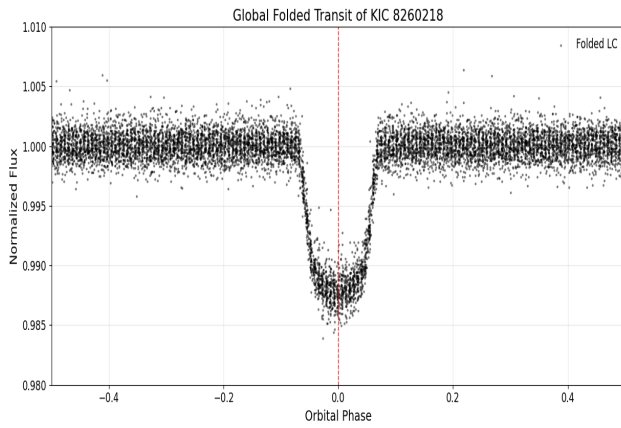
(a) Global view of KIC 6300348
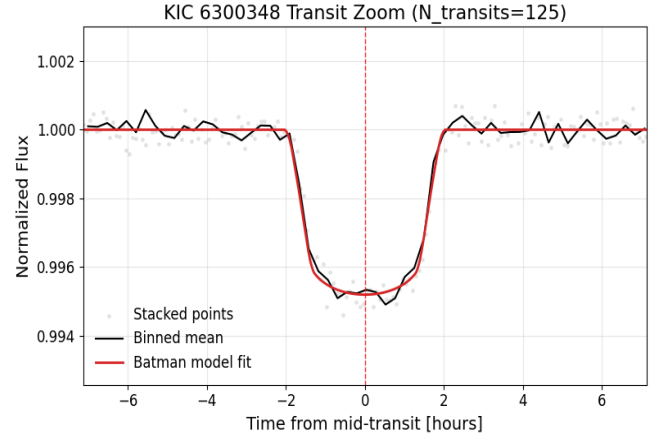
(b) Zoomed transit view of KIC 6300348

Figure 11. Phase-folded light curve of KIC 6300348. Left: full orbital phase. Right: zoomed transit morphology.
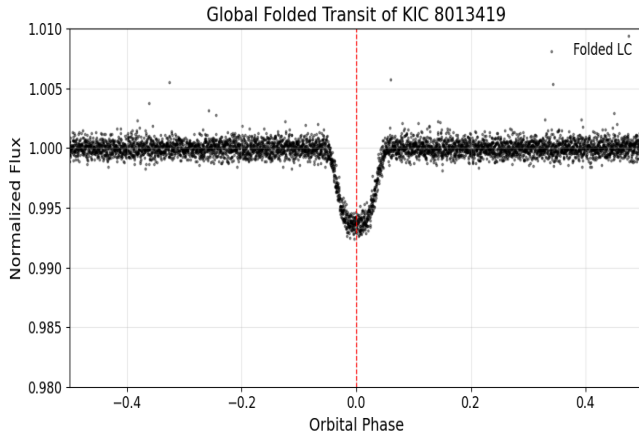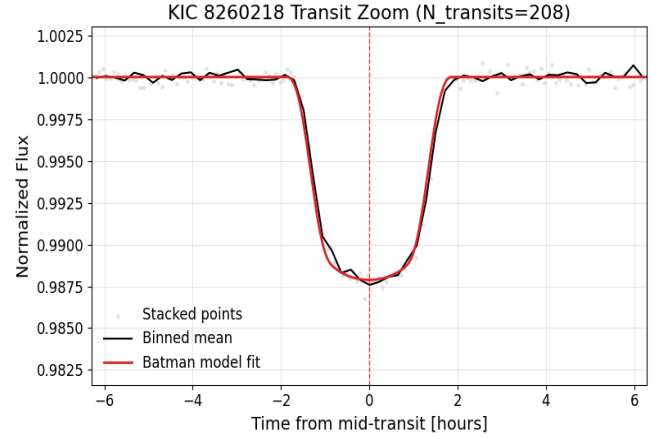

(a) Global view of KIC 8012419

(b) Zoomed transit view of KIC 8012419

Figure 12. Phase-folded light curve of KIC 8012419.

(a) Global view of KIC 8260218

(b) Zoomed transit view of KIC 8260218

Figure 13. Phase-folded light curve of KIC 8260218.